

ECE4574 – Large-Scale SW Development for Engineering Systems

Lecture 6 – Review of Networks

Creed Jones, PhD

Course Updates

- Your project proposal is due TODAY, 11:59 PM
 - It's a group assignment – only one person (the Scrum Master) need submit it
- Homework 1 is due Friday, September 22
- Quiz 2 is TODAY, September 11
 - 7 PM to 1 AM Eastern time
 - 20 minute time limit
 - open notes, no help from anyone else
 - covers lectures 3 and 4
- No office hours on Tuesday, Sept 12
 - schedule conflict

Today's Objectives

Brief review of networking

- The layered model
- Application layer
 - HTTP
 - port numbers
- Transport layer
- Network layer
- Link and physical layers

A Network is a connection between computers that can be used to exchange data; network interactions are defined in *protocols*

Standard – not custom

Packet-based – not byte by byte

Multi-use – not dedicated to one app

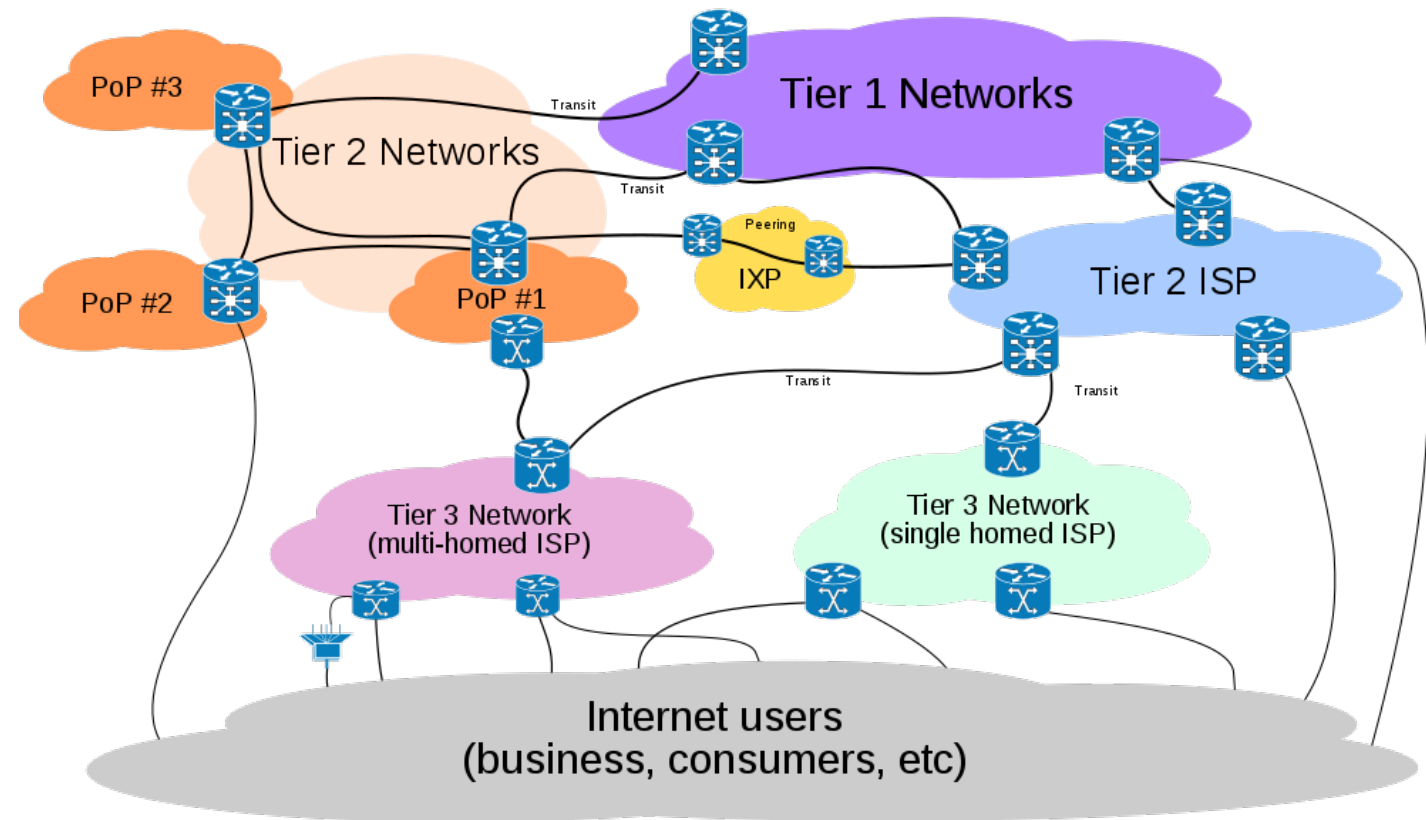
Multiple connections – not one-to-one

- A protocol is set of rules governing the exchange of data between two entities
 - used for communication between entities that can exchange information
 - for two entities to communicate successfully they must “speak the same language”
- Key elements of a protocol are:

Syntax	Semantics	Timing
<ul style="list-style-type: none">• includes such things as data format and signal levels	<ul style="list-style-type: none">• includes control information for coordination and error handling	<ul style="list-style-type: none">• includes speed matching and sequencing

A network can be divided into systems at the edge and the core; the core exists to provide connections between nodes on the edge

- This is the ***end-to-end network principle***
 - The application is characterized by what happens in system on the edge
 - Core systems just support message passing
- **network edge:** applications and hosts
- **network core:**
 - routers
 - gateways
 - network of networks
- **access networks, physical media:** communication links



We can connect edge nodes in several ways: for example, TCP is a connection-oriented service

Goal: data transfer between end systems

- *handshaking*: setup (prepare for) data transfer ahead of time
 - Hello, hello back human protocol
 - *set up "state"* in two communicating hosts
- TCP - Transmission Control Protocol
 - Internet's connection-oriented service

TCP service [RFC 793]

- *reliable, in-order* byte-stream data transfer
 - loss: acknowledgements and retransmissions
- *flow control*:
 - sender won't overwhelm receiver
- *congestion control*:
 - senders "slow down sending rate" when network congested

While TCP is based on establishing and using a connection, UDP is a connectionless service

Goal: data transfer between end systems

- same as before!
- **UDP** - User Datagram Protocol [RFC 768]:
 - connectionless
 - unreliable data transfer
 - no flow control
 - no congestion control

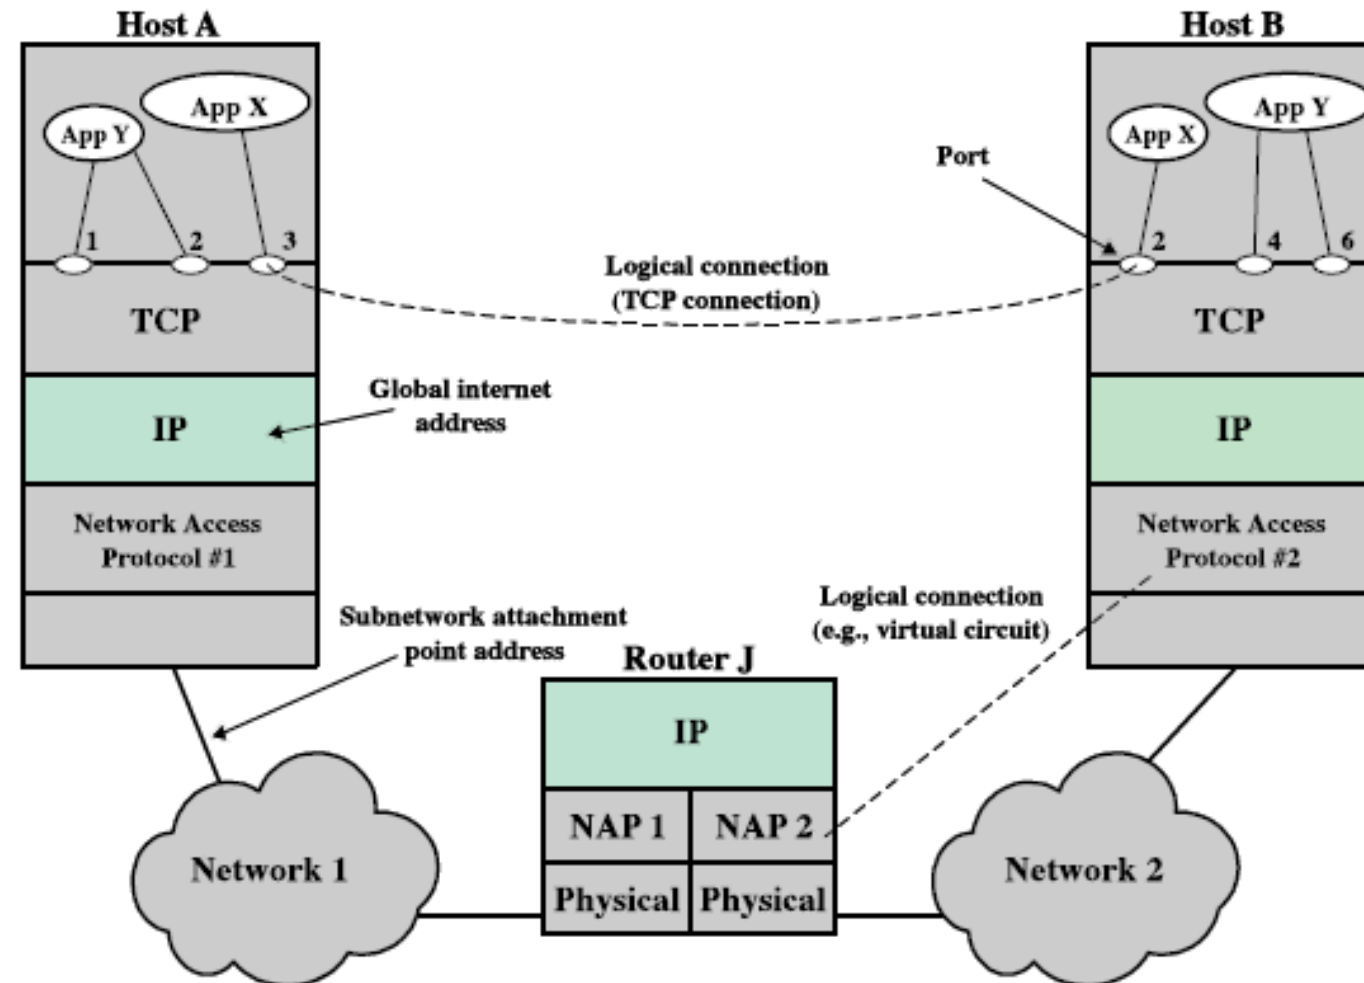
Apps using TCP:

- HTTP/HTTPS (Web), FTP (file transfer), Telnet (remote login), SMTP/IMAP (email)

Apps using UDP:

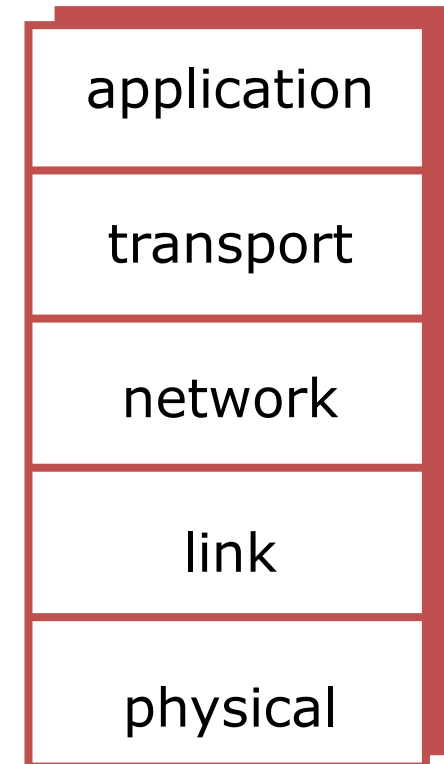
- streaming media, teleconferencing, DNS, Internet telephony

Network Protocols are Layered; each layer handles certain tasks and connects with layers above and below

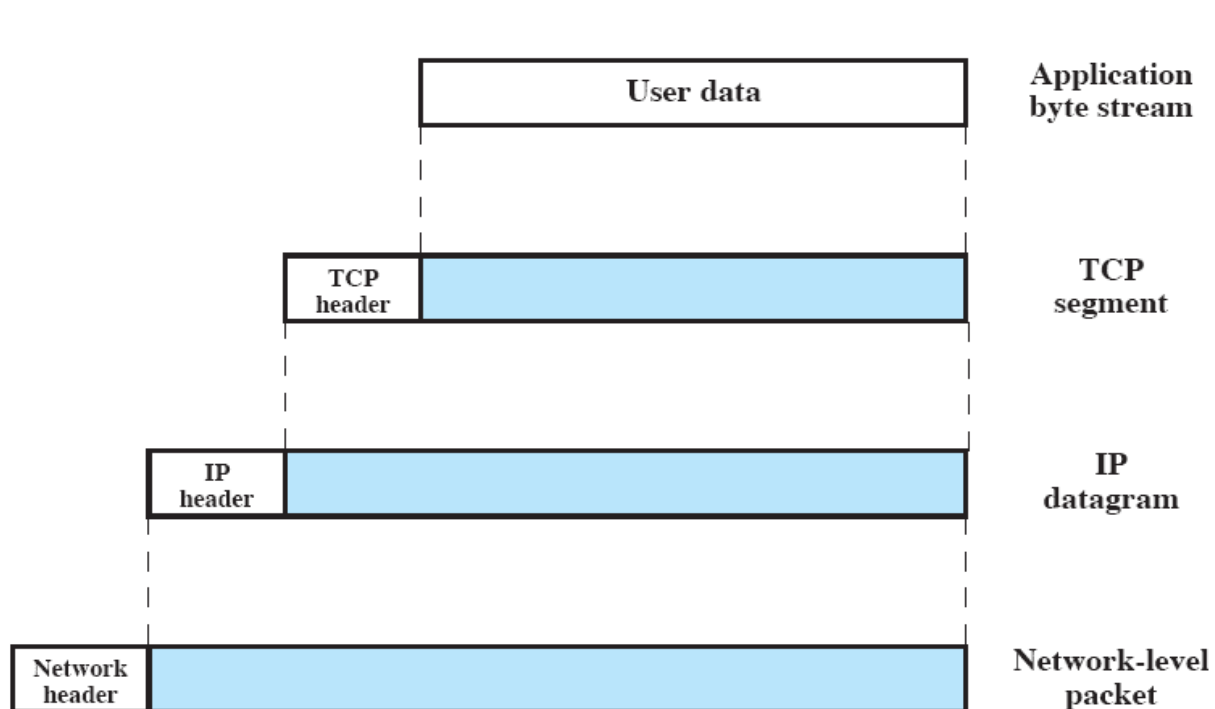


The Internet protocol stack consists of layers

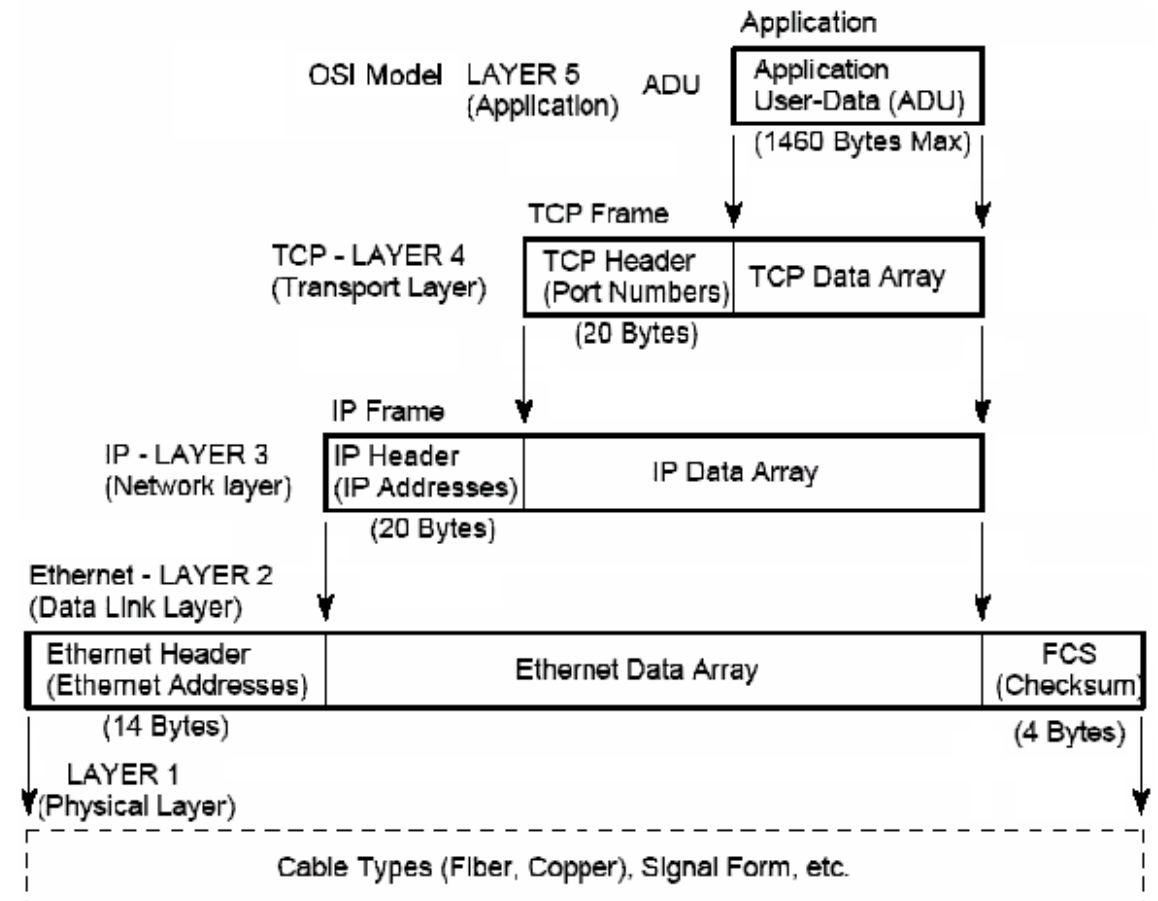
- **application:** supporting network applications
 - FTP, SMTP, HTTP
- **transport:** host-host data transfer
 - TCP, UDP
- **network:** routing of datagrams from source to destination
 - IP, routing protocols
- **link:** data transfer between neighboring network elements
 - PPP, Ethernet
- **physical:** bits on a cable, on Wifi or cell network

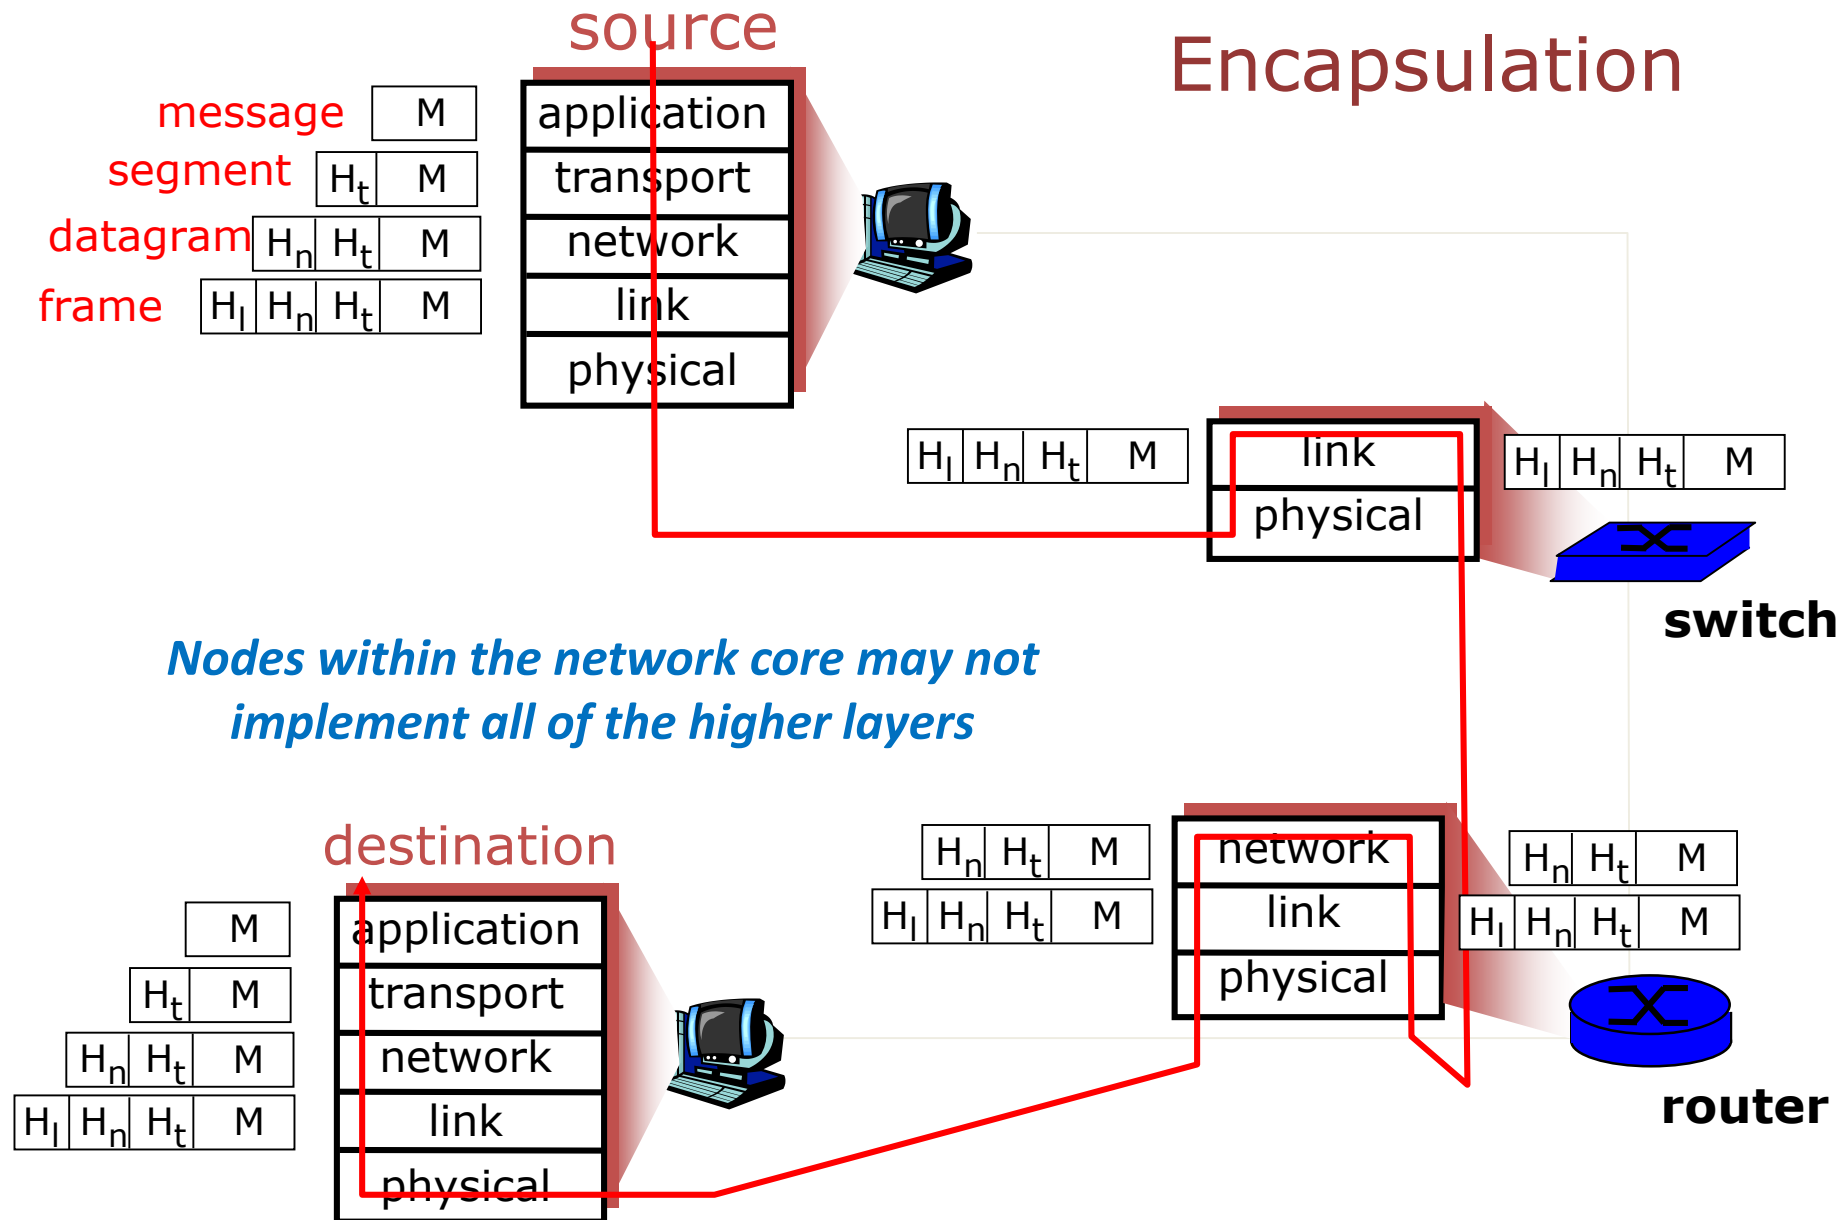


Data is encapsulated by lower layers

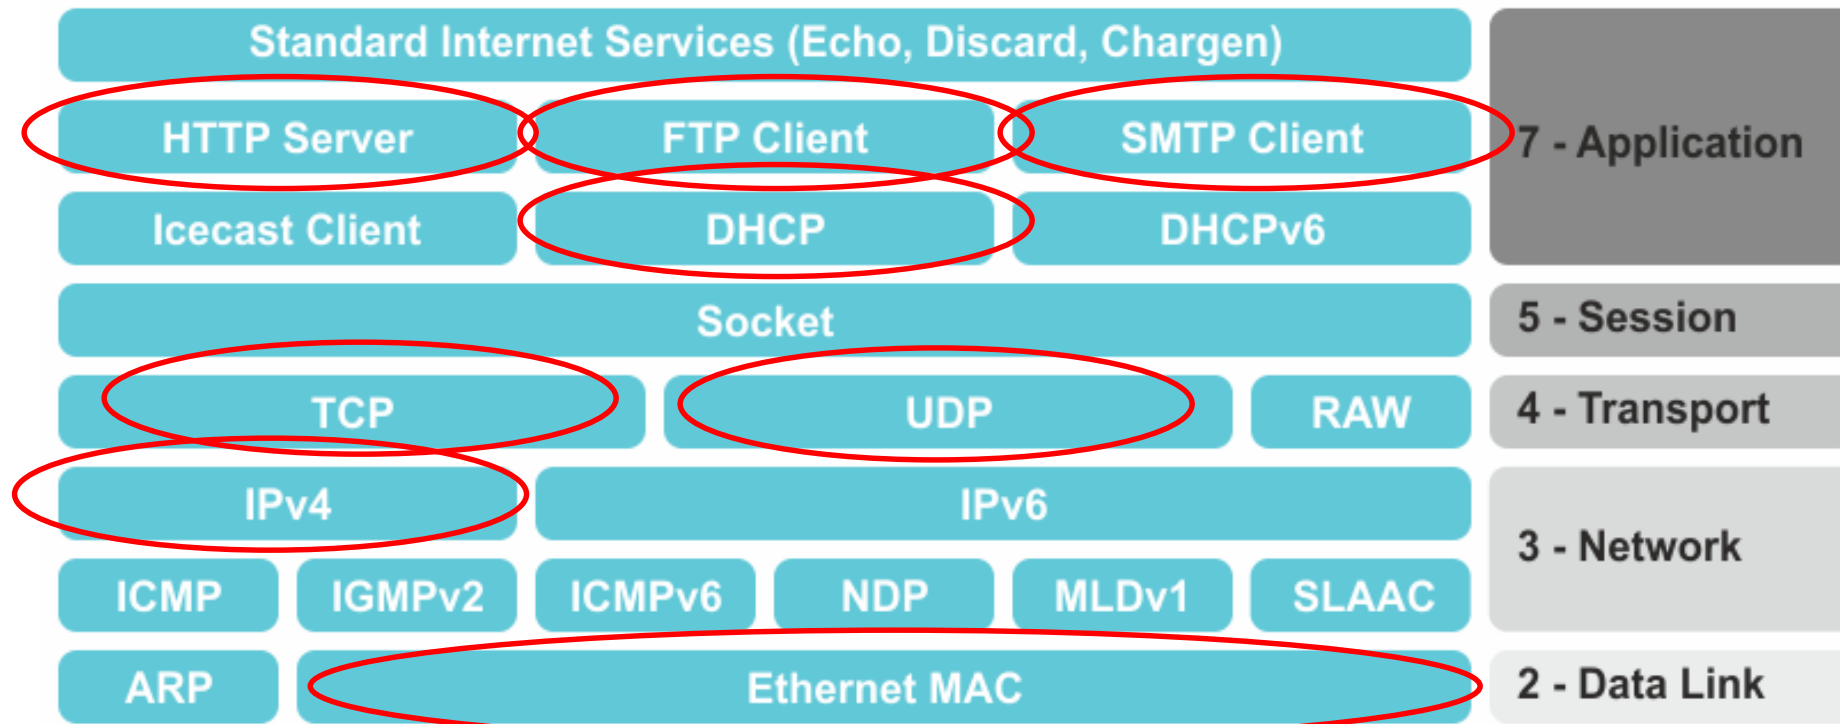


CONSTRUCTION OF A TCP/IP-ETHERNET DATA PACKET





Many protocols exist at most levels of the TCP/IP
 Protocol stack:
 these are the most important ones



Internet protocols are defined in documents called RFCs (Request for Comments)

- These standard definitions of the protocols are not laid down by some authority, in most cases
- They were developed and are followed by common consent
- Initial RFCs were issued and changes made
 - When the changes settled down, the RFC was considered the defining document
- You can see them at <http://www.rfc-editor.org/rfc.html>
 - Search for "HTTP" – the key docs are RFC 1945, RFC 7230 and RFC 2660

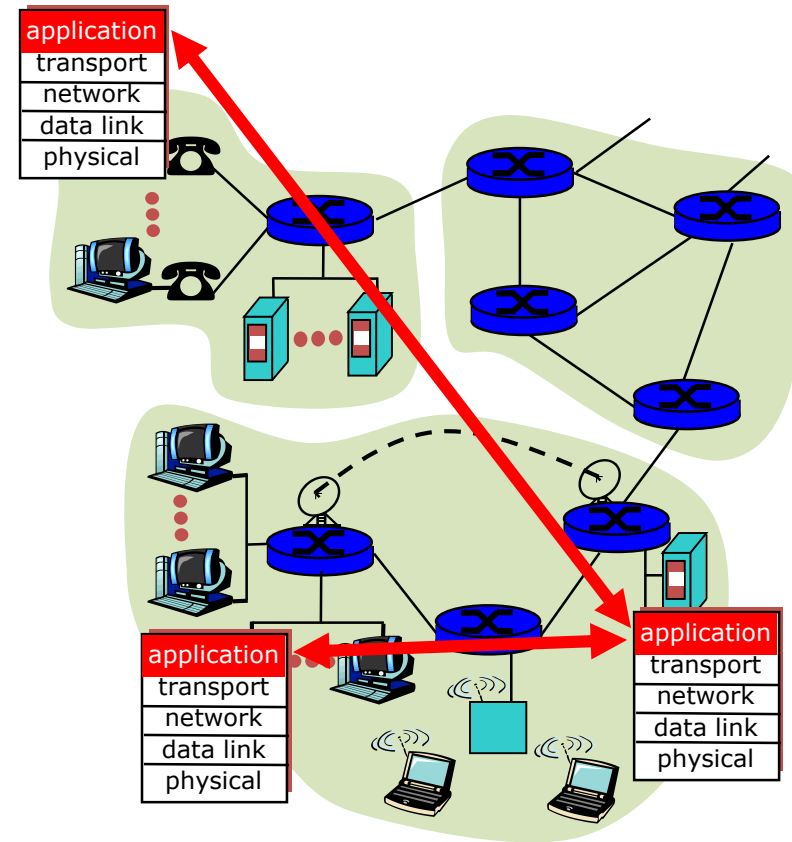
Creating a network application

Write programs that

- run on different end systems and
- communicate over a network.
- e.g., Web: Web server software communicates with browser software

Less software is written for devices in network core

- network core devices do not run user application code
- application on end systems allows for rapid app development, propagation



APPLICATION LAYER

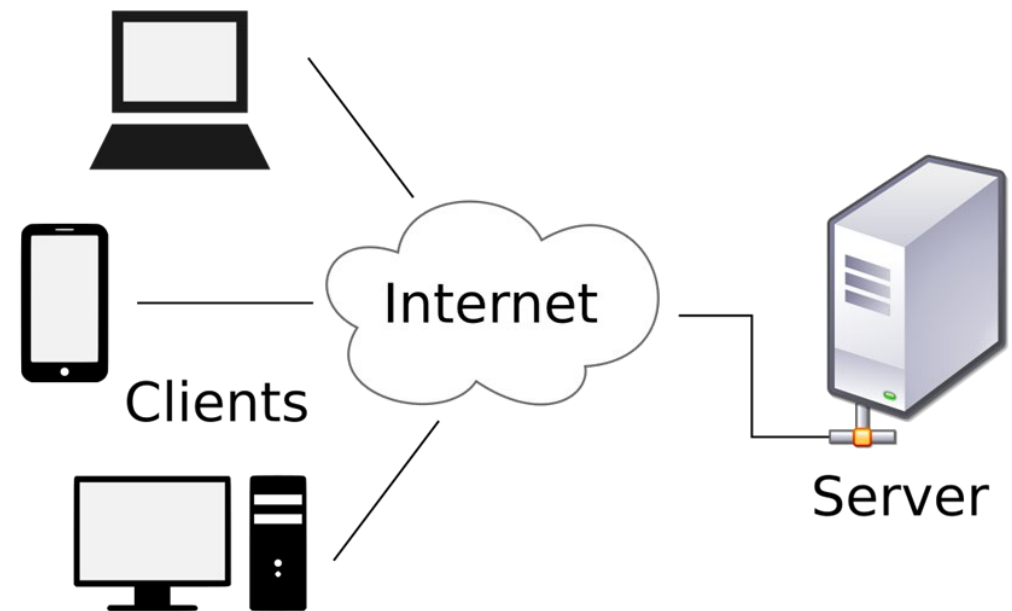
Client-server architecture

server:

- always-on host
- often, permanent IP address
- server farms for scaling

clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other



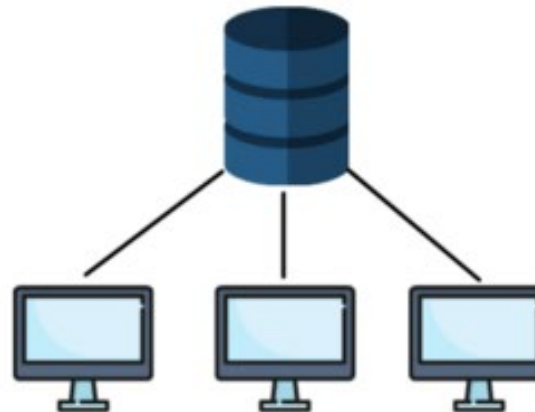
Pure Peer to Peer (P2P) architecture

- no always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses
- examples: BitTorrent and Akamai

Highly scalable

But difficult to manage

Client Server Architecture



Peer to Peer Architecture



Hybrid of client-server and P2P

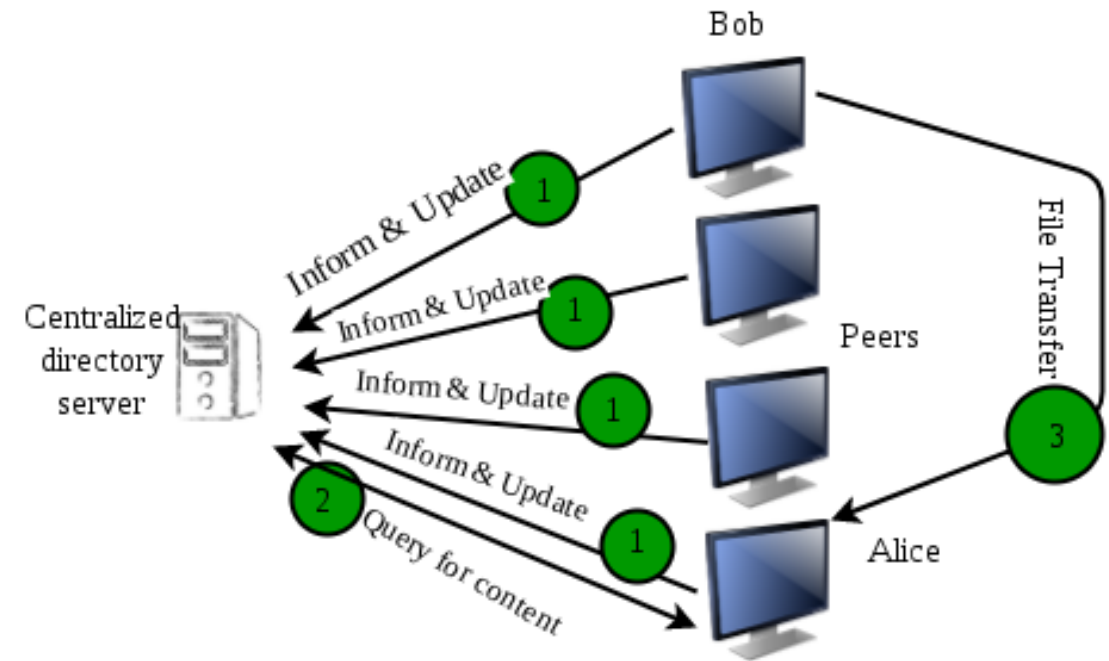
Instant messaging

- Chatting between two users can be P2P
- Presence detection/location centralized:
 - User registers its IP address with central server when it comes online
 - User contacts central server to find IP addresses of buddies

File sharing

- Linux distribution

Bitcoin



P2P paradigm with a centralised directory

Processes communicating

Process: program running within a host.

- within same host, two processes communicate using **inter-process communication** (defined by OS).
- processes in different hosts communicate by exchanging **messages**

Client process: process that initiates communication

Server process: process that waits to be contacted

- Note: applications with P2P architectures have client processes & server processes

Addressing processes uses both the IP address and the *port number*

- Identifier includes both the IP address and **port numbers** associated with the process on the host.
- Example port numbers:
 - HTTP server: 80
 - Mail server: 25
- In programs, a network *socket* is opened on a particular address and port number
- For a process to receive messages, it must have an identifier
- A host has a unique 32-bit IP address
- **Q:** does the IP address of the host on which the process runs suffice for identifying the process?
- **Answer:** No, many processes can be running on same host

PORT NUMBERS

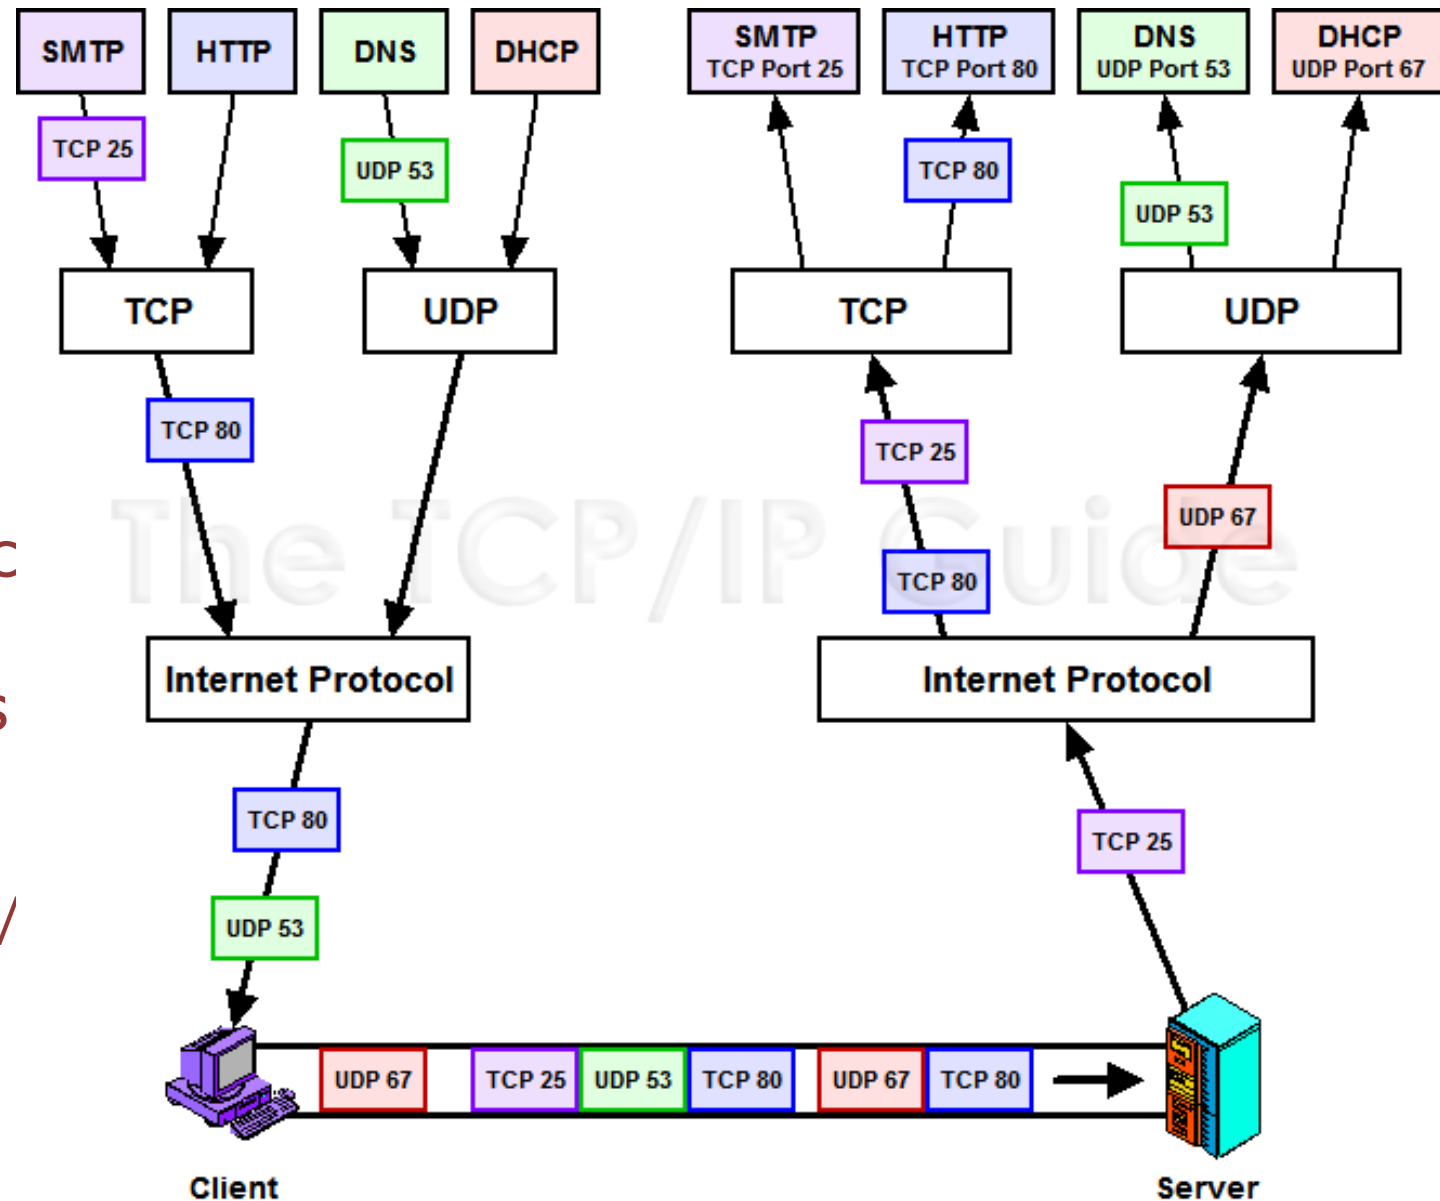
Recall that port numbers are used to route incoming network traffic to the right application

- The IP address (derived from the URL using the Domain Name Service, DNS) will get traffic to the right host
- On that host, different applications are looking for traffic that carries the proper port number
- More specifically, a network *socket* is created by the application code on the destination, and is associated with a single port number and destination IP address
- Many port numbers are universally connected with certain types of traffic
 - these are called the well-known port numbers

Port numbers are
 processed by the
 transport layer (the IP
 implementation) and
 allow direction of traffic
 to and from different
 processes on the hosts

- from

<http://www.tcpipguide.com/>



The “well-known” port numbers are universally agreed to route traffic to specific applications – NOTE that the application can still split data up by other means!

PORT NUMBER	TRANSPORT PROTOCOL	SERVICE NAME	RFC
20, 21	TCP	File Transfer Protocol (FTP)	RFC 959
22	TCP and UDP	Secure Shell (SSH)	RFC 4250-4256
23	TCP	Telnet	RFC 854
25	TCP	Simple Mail Transfer Protocol (SMTP)	RFC 5321
53	TCP and UDP	Domain Name Server (DNS)	RFC 1034-1035
67, 68	UDP	Dynamic Host Configuration Protocol (DHCP)	RFC 2131
69	UDP	Trivial File Transfer Protocol (TFTP)	RFC 1350
80	TCP	HyperText Transfer Protocol (HTTP)	RFC 2616
110	TCP	Post Office Protocol (POP3)	RFC 1939
119	TCP	Network News Transport Protocol (NNTP)	RFC 8977
123	UDP	Network Time Protocol (NTP)	RFC 5905
135-139	TCP and UDP	NetBIOS	RFC 1001-1002
143	TCP and UDP	Internet Message Access Protocol (IMAP4)	RFC 3501
161, 162	TCP and UDP	Simple Network Management Protocol (SNMP)	RFC 1901-1908, 3411-3418
179	TCP	Border Gateway Protocol (BGP)	RFC 4271
389	TCP and UDP	Lightweight Directory Access Protocol	RFC 4510
443	TCP and UDP	HTTP with Secure Sockets Layer (SSL)	RFC 2818
500	UDP	Internet Security Association and Key Management Protocol (ISAKMP) / Internet Key Exchange (IKE)	RFC 2408 - 2409
636	TCP and UDP	Lightweight Directory Access Protocol over TLS/SSL (LDAPS)	RFC 4513
989/990	TCP	FTP over TLS/SSL	RFC 4217

<https://ipwithease.com>

7 Echo	554 RTSP	2745 Bagle.H	6891-6901 Windows Live
19 Chargen	546-547 DHCPv6	2967 Symantec AV	6970 Quicktime
20-21 FTP	560 rmonitor	3050 Interbase DB	7212 GhostSurf
22 SSH/SCP	563 NNTP over SSL	3074 XBOX Live	7648-7649 CU-SeeMe
23 Telnet	587 SMTP	3124 HTTP Proxy	8000 Internet Radio
25 SMTP	591 FileMaker	3127 MyDoom	8080 HTTP Proxy
42 WINS Replication	593 Microsoft DCOM	3128 HTTP Proxy	8086-8087 Kaspersky AV
43 WHOIS	631 Internet Printing	3222 GLBP	8118 Privoxy
49 TACACS	636 LDAP over SSL	3260 iSCSI Target	8200 VMware Server
53 DNS	639 MSDP (PIM)	3306 MySQL	8500 Adobe ColdFusion
67-68 DHCP/BOOTP	646 LDP (MPLS)	3389 Terminal Server	8767 TeamSpeak
69 TFTP	691 MS Exchange	3689 iTunes	8866 Bagle.B
70 Gopher	860 iSCSI	3690 Subversion	9100 HP JetDirect
79 Finger	873 rsync	3724 World of Warcraft	9101-9103 Bacula
80 HTTP	902 VMware Server	3784-3785 Ventrilo	9119 MXit
88 Kerberos	989-990 FTP over SSL	4333 mSQL	9800 WebDAV
102 MS Exchange	993 IMAP4 over SSL	4444 Blaster	9898 Dabber
110 POP3	995 POP3 over SSL	4664 Google Desktop	9988 Rbot/Spybot
113 Ident	1025 Microsoft RPC	4672 eMule	9999 Urchin
119 NNTP (Usenet)	1026-1029 Windows Messenger	4899 Radmin	10000 Webmin
123 NTP	1080 SOCKS Proxy	5000 UPnP	10000 BackupExec
135 Microsoft RPC	1080 MyDoom	5001 Slingbox	10113-10116 NetIQ
137-139 NetBIOS	1194 OpenVPN	5001 iperf	11371 OpenPGP
143 IMAP4	1214 Kazaa	5004-5005 RTP	12035-12036 Second Life
161-162 SNMP	1241 Nessus	5050 Yahoo! Messenger	12345 NetBus
177 XDMCP	1311 Dell OpenManage	5060 SIP	13720-13721 NetBackup
179 BGP	1337 WASTE	5190 AIM/ICQ	14567 Battlefield
201 AppleTalk	1433-1434 Microsoft SQL	5222-5223 XMPP/Jabber	15118 Dipnet/Oddbob
264 BGMP	1512 WINS	5432 PostgreSQL	19226 AdminSecure
318 TSP	1589 Cisco VQP	5500 VNC Server	19638 Ensism
381-383 HP Openview	1701 L2TP	5554 Sasser	20000 Usermin
389 LDAP	1723 MS PPTP	5631-5632 pcAnywhere	24800 Synergy
411-412 Direct Connect	1725 Steam	5800 VNC over HTTP	25999 Xfire
443 HTTP over SSL	1741 CiscoWorks 2000	5900+ VNC Server	27015 Half-Life
445 Microsoft DS	1755 MS Media Server	6000-6001 X11	27374 Sub7
464 Kerberos	1812-1813 RADIUS	6112 Battle.net	28960 Call of Duty
465 SMTP over SSL	1863 MSN	6129 DameWare	31337 Back Orifice
497 Retrospect	1985 Cisco HSRP	6257 WinMX	33434+ traceroute
500 ISAKMP	2000 Cisco SCCP	6346-6347 Gnutella	
512 rexec	2002 Cisco ACS	6500 GameSpy Arcade	
513 rlogin	2049 NFS	6566 SANE	
514 syslog	2082-2083 cPanel	6588 AnalogX	
515 LPD/LPR	2100 Oracle XDB	6665-6669 IRC	
520 RIP	2222 DirectAdmin	6679/6697 IRC over SSL	
521 RIPng (IPv6)	2302 Halo	6699 Napster	
540 UUCP	2483-2484 Oracle DB	6881-6999 BitTorrent	

Legends
Chat
Encrypted
Gaming
Malicious
Peer to Peer
Streaming

Here is a partial list of port numbers – note that these include:

- protocols
- applications
- OS utilities
- comm (VOIP)
- games

App-layer protocol defines

- Types of messages exchanged, e.g., request & response messages
- Syntax of message types: what fields in messages & how fields are delineated
- Semantics of the fields, i.e., meaning of information in fields
- Rules for when and how processes send & respond to messages

Public-domain protocols:

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP

Proprietary protocols:

- e.g., Skype

What transport service does an app need?

Data loss

- some apps (e.g., audio) can tolerate some loss
- other apps (e.g., file transfer, telnet) require 100% reliable data transfer

Bandwidth

- some apps (e.g., multimedia) require minimum amount of bandwidth to be “effective”
- other apps (“elastic apps”) make use of whatever bandwidth they get

Timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

Transport service requirements of common apps

Application	Data loss	Bandwidth	Time Sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100's msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100's msec
instant messaging	no loss	elastic	yes and no

Internet transport protocols services

TCP service:

- *connection-oriented*: setup required between client and server processes
- *reliable transport* between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum bandwidth guarantees

UDP service:

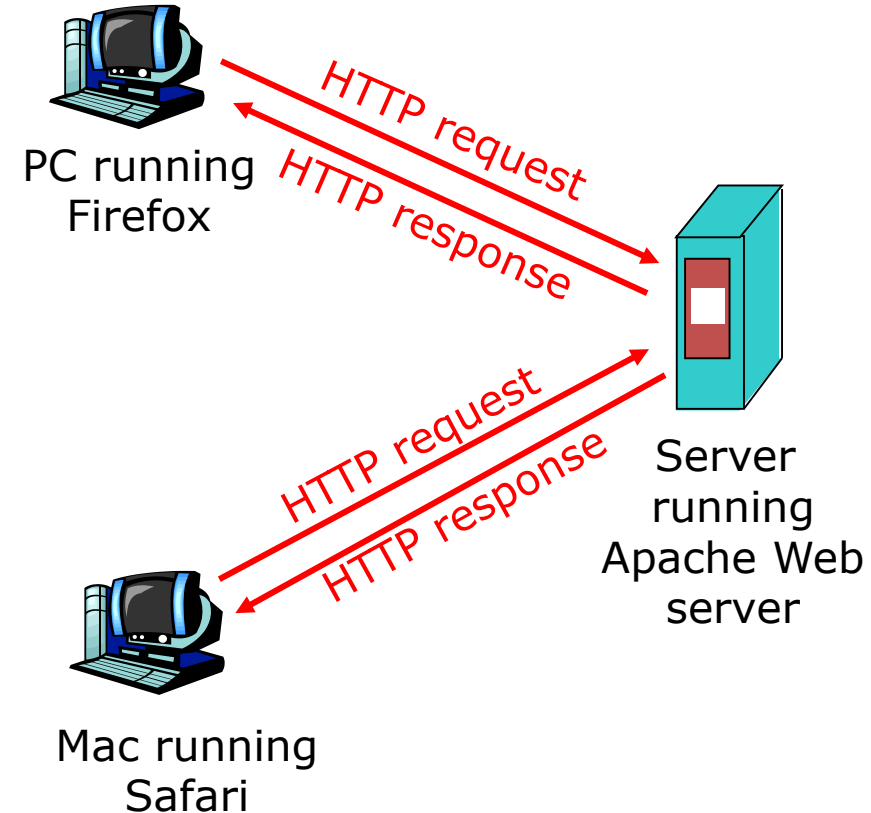
- unreliable data transfer between sending and receiving process
- does not provide: connection setup, reliability, flow control, congestion control, timing, or bandwidth guarantee

Internet apps: application, transport protocols

Application	Application layer protocol	Underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP [RFC 793]
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 7540]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	proprietary (e.g. RealNetworks)	TCP or UDP
Internet telephony	proprietary (e.g., Vonage, Dialpad)	typically UDP

HTTP – the hypertext transfer protocol

- Web's application layer protocol
- client/server model
 - *client*: browser that requests, receives, "displays" Web objects
 - *server*: Web server sends objects in response to requests
- HTTP 1.1: [RFC 2330](#)
- HTTP 2.0: [RFC 7540](#)



HTTP overview (continued)

Uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is "stateless"

- server maintains no information about past client requests

Protocols that maintain "state" are complex!

- past history (state) must be maintained
- if server/client crashes, their views of "state" may be inconsistent, must be reconciled

HTTP request message

- two types of HTTP messages: *request, response*
- HTTP request message:
 - ASCII text (human-readable format)

request line
(GET, POST,
HEAD commands)

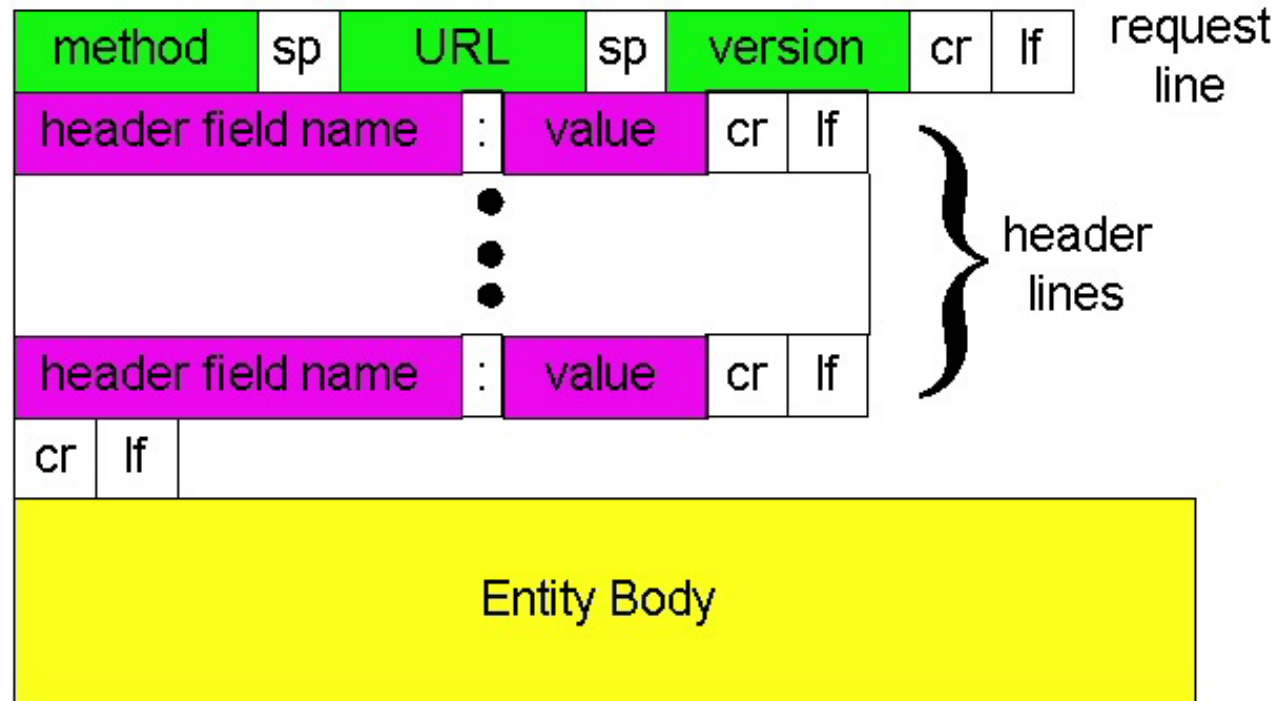
header
lines

```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
```

Carriage return,
line feed
indicates end
of message

(extra carriage return, line feed)

HTTP request message: general format



```
GET /somedir/page.html HTTP/1.1
Host: www.someschool.edu
User-agent: Mozilla/4.0
Connection: close
Accept-language: fr
```

Clients can specify various operations, and provide data to servers, in several ways

HTTP supports several message types

- GET, POST, HEAD
- PUT
 - uploads file in entity body to path specified in URL field
- DELETE
 - deletes file specified in the URL field

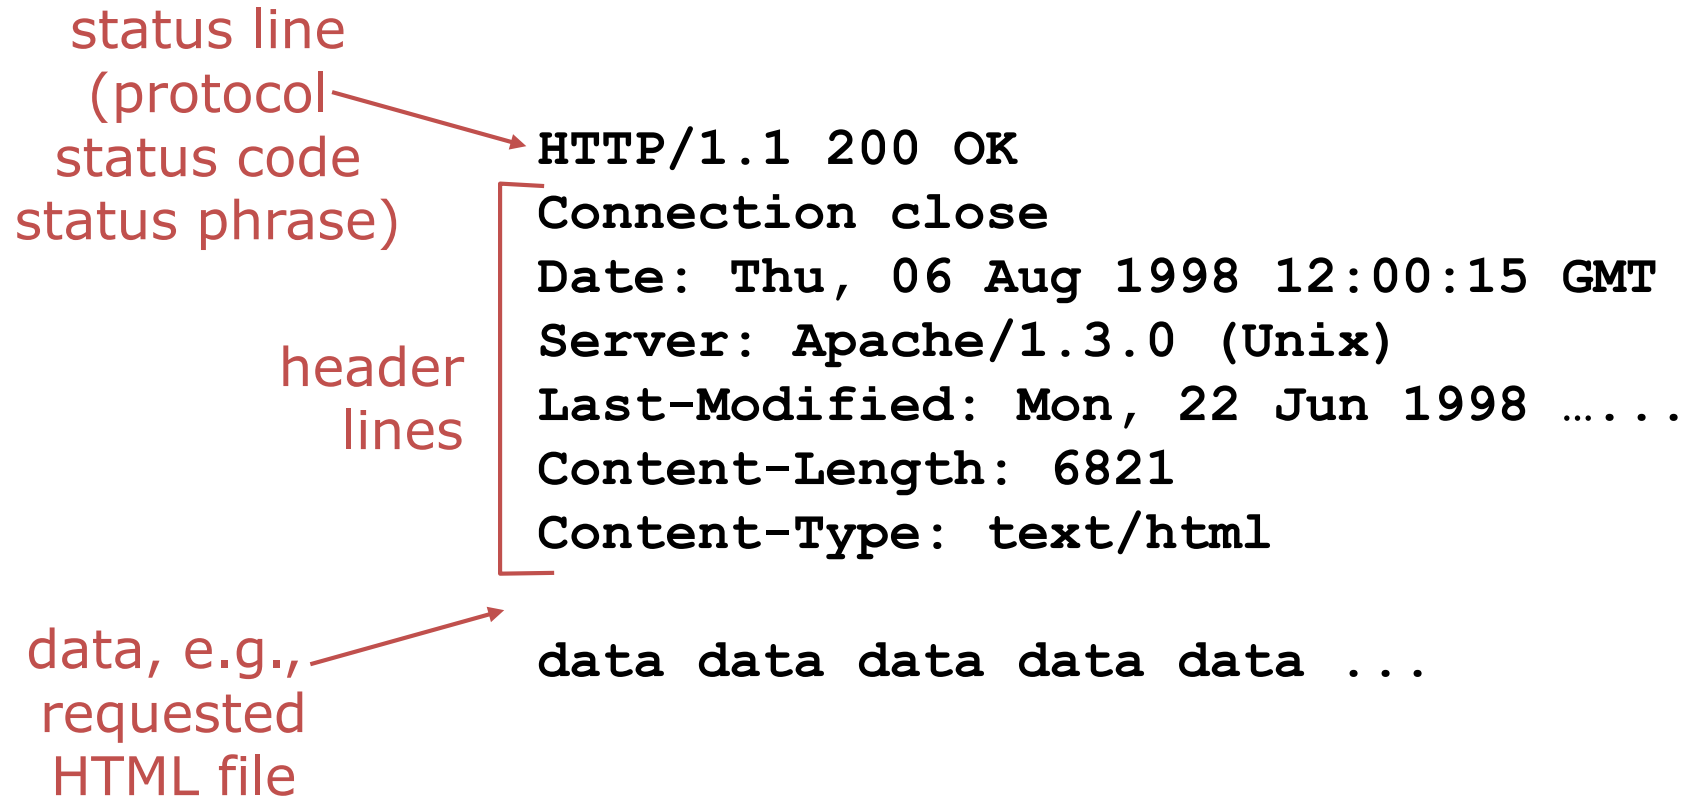
Post method of data upload:

- Web page often includes form input
- Input is uploaded to server in entity body

URL method of data upload:

- Uses GET method
- Input is uploaded in URL field of request line:
- `www.somesite.com/animalsearch?monkeys&banana`

HTTP response message



HTTP response status codes

In first line in server->client response message.

A few sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (Location:)

400 Bad Request

- request message not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

HTTP Status Codes

Level 200 (Success)

200 : OK

201 : Created

203 : Non-Authoritative
Information

204 : No Content

Level 400

400 : Bad Request

401 : Unauthorized

403 : Forbidden

404 : Not Found

409 : Conflict

Level 500

500 : Internal Server Error

503 : Service Unavailable

501 : Not Implemented

504 : Gateway Timeout

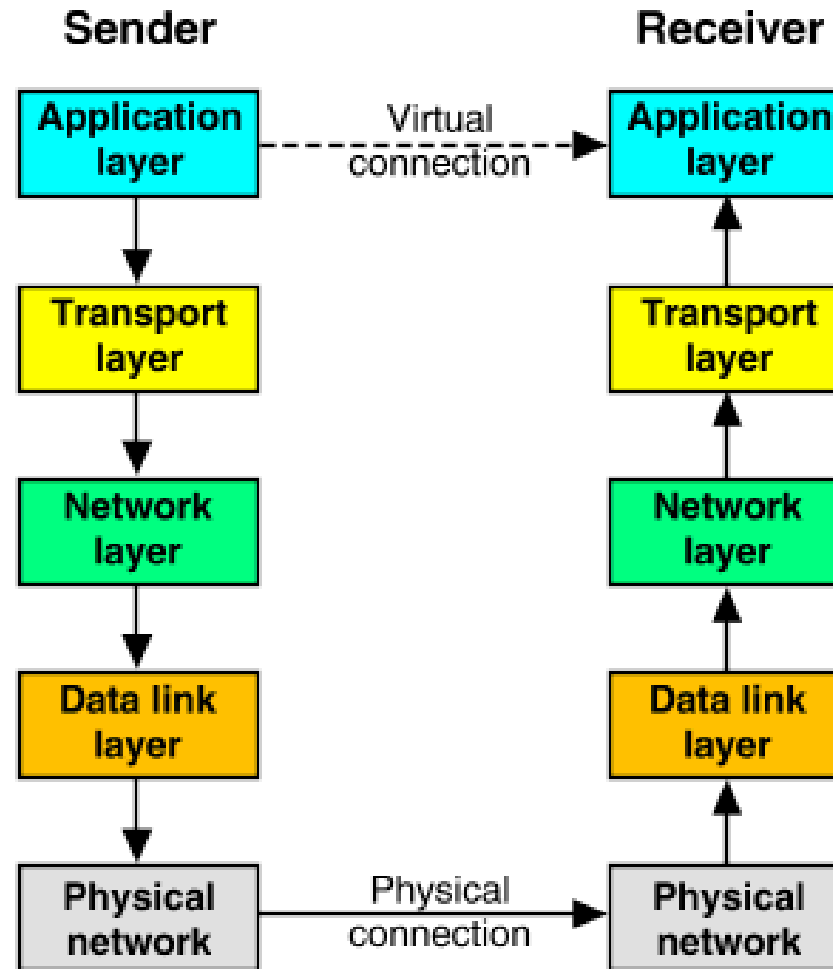
599 : Network timeout

502 : Bad Gateway

THE TRANSPORT LAYER

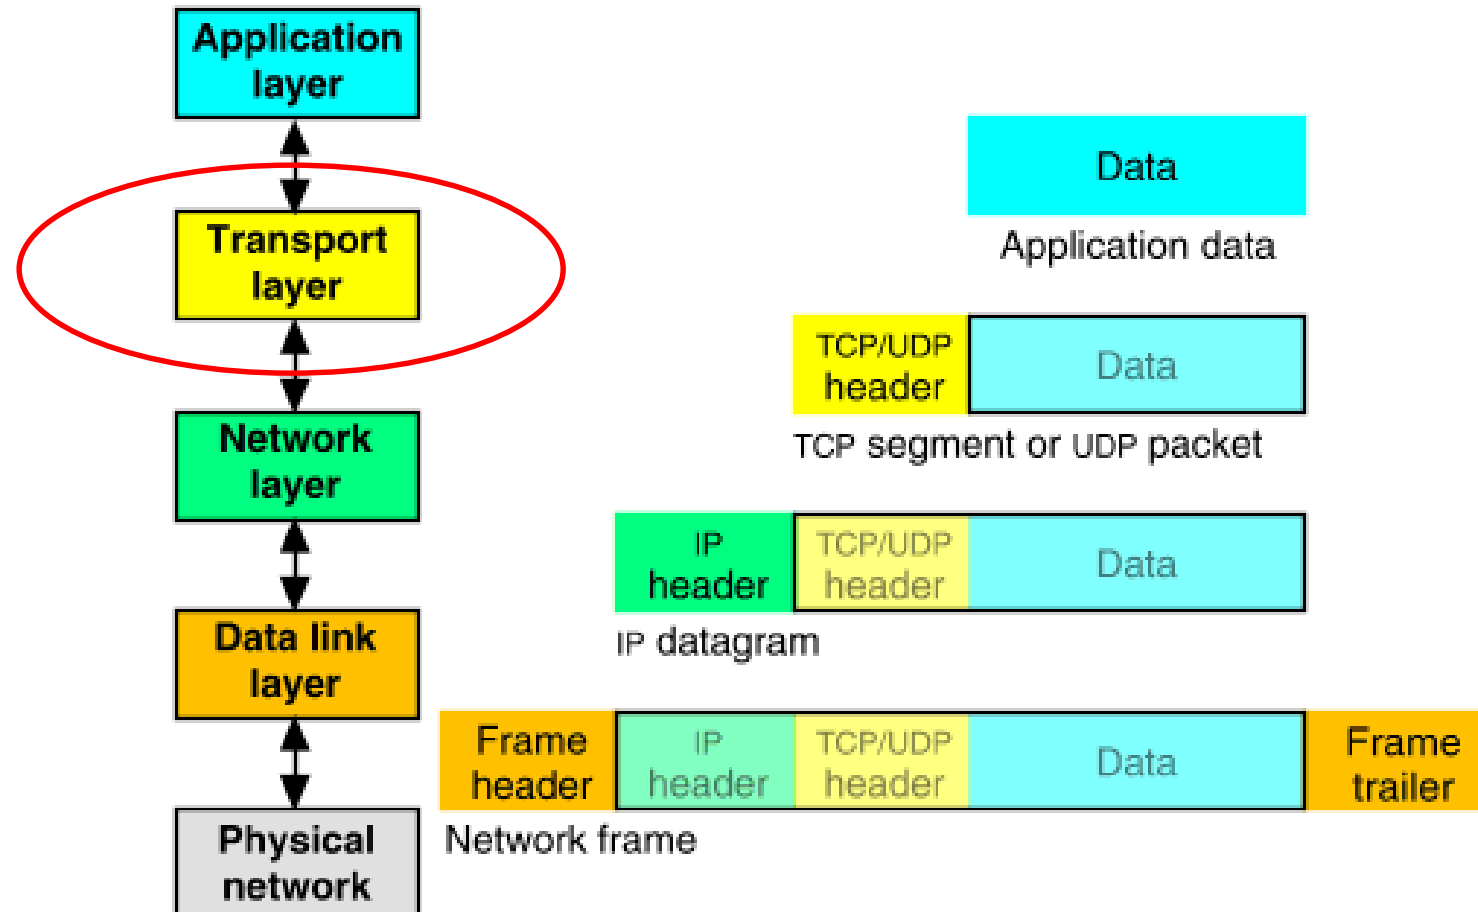
TCP/IP protocol stack

- Recall the nature of the TCP/IP network stack
 - An example of any layered set of protocols
- Each layer consists of a protocol
 - a set of messages and data formats
- Data is encapsulated as it passes between layers



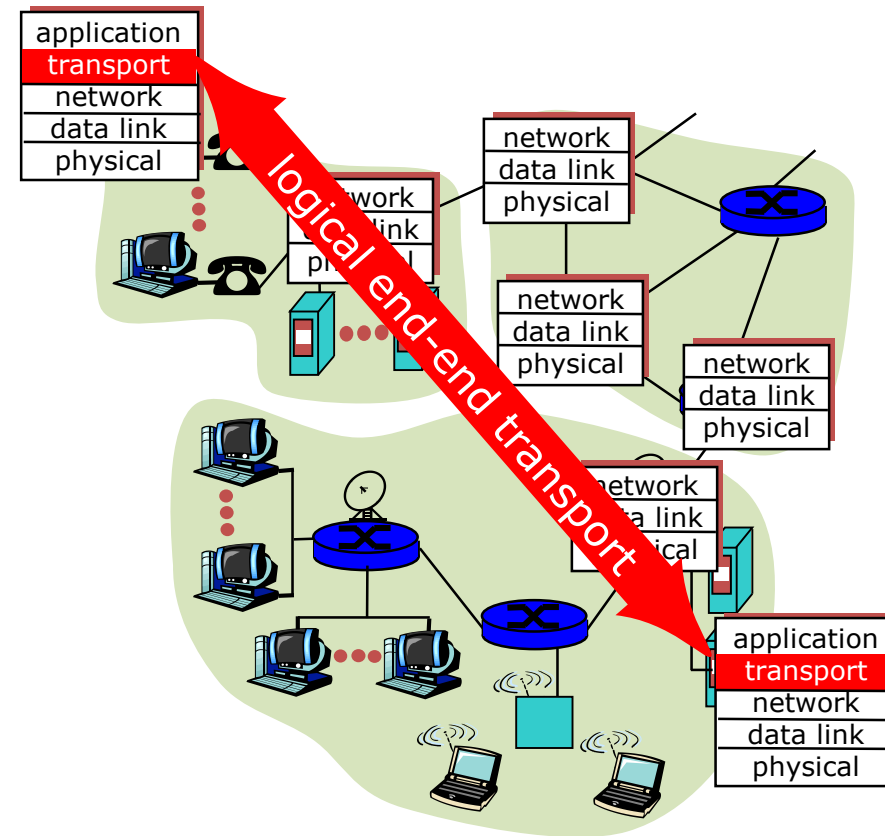
from http://uw713doc.sco.com/en/NET_tcpip/graphics/encapsulation.gif

TCP or UDP work at the transport layer



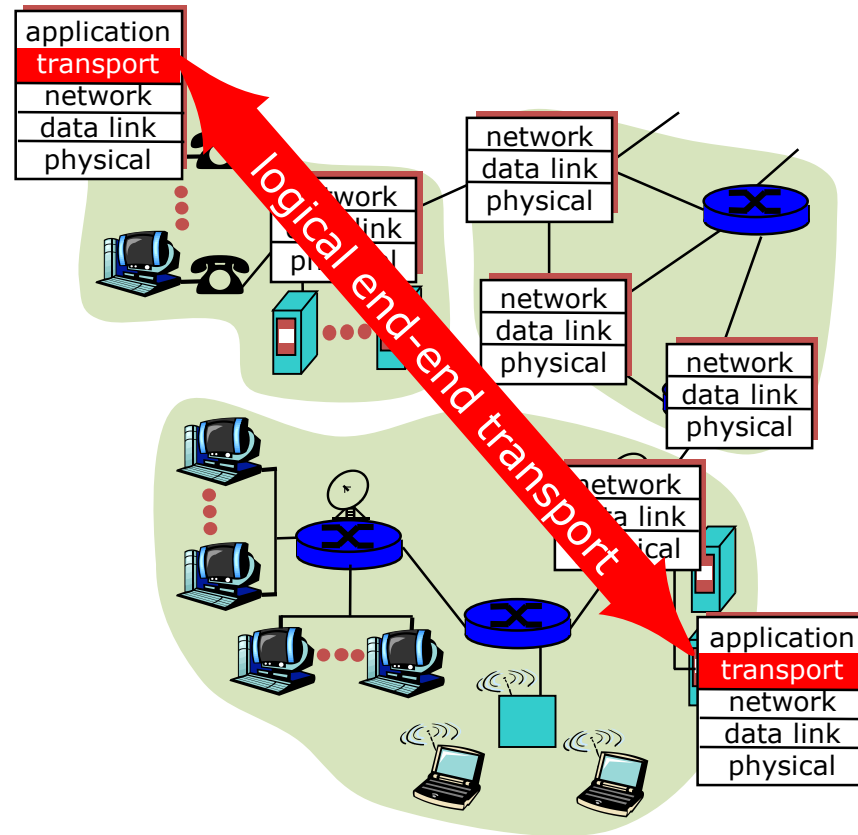
Transport services and protocols

- provide *logical communication* between app processes running on different hosts
- transport protocols run in end systems
 - send side: breaks app messages into *segments*, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- more than one transport protocol available to apps
 - Internet: TCP and UDP



Internet transport-layer protocols

- reliable, in-order delivery (TCP)
 - congestion control
 - flow control
 - connection setup
- unreliable, unordered delivery: UDP
 - no-frills extension of “best-effort” IP
- services not available:
 - delay guarantees
 - bandwidth guarantees



UDP: User Datagram Protocol [RFC 768]

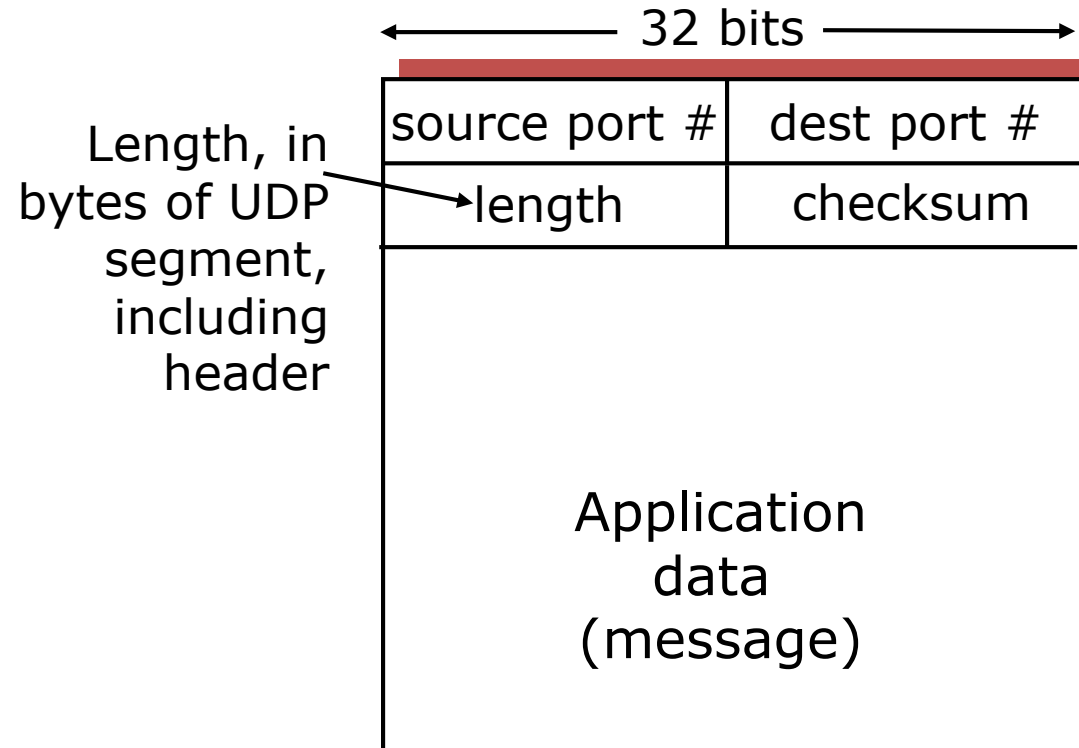
- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
 - lost
 - delivered out of order to app
- *connectionless*:
 - no handshaking between UDP sender, receiver
 - each UDP segment handled independently of others

Why is there a UDP?

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small segment header
- no congestion control: UDP can blast away as fast as desired

User Datagram Protocol

- often used for streaming multimedia apps
 - loss tolerant
 - rate sensitive
- other UDP uses
 - DNS
 - SNMP
- reliable transfer over UDP: add reliability at application layer
 - application-specific error recovery!



UDP segment format

UDP checksum

Goal: detect “errors” (e.g., flipped bits) in transmitted segment

Sender:

- treat segment contents as sequence of 16-bit integers
- checksum: addition (1’s complement sum) of segment contents
- sender puts checksum value into UDP checksum field

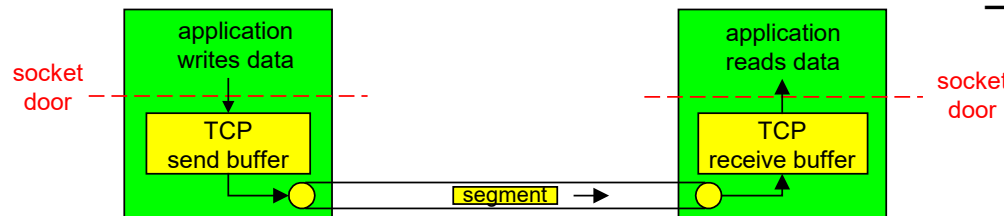
Receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected. *But maybe errors nonetheless?*
More later

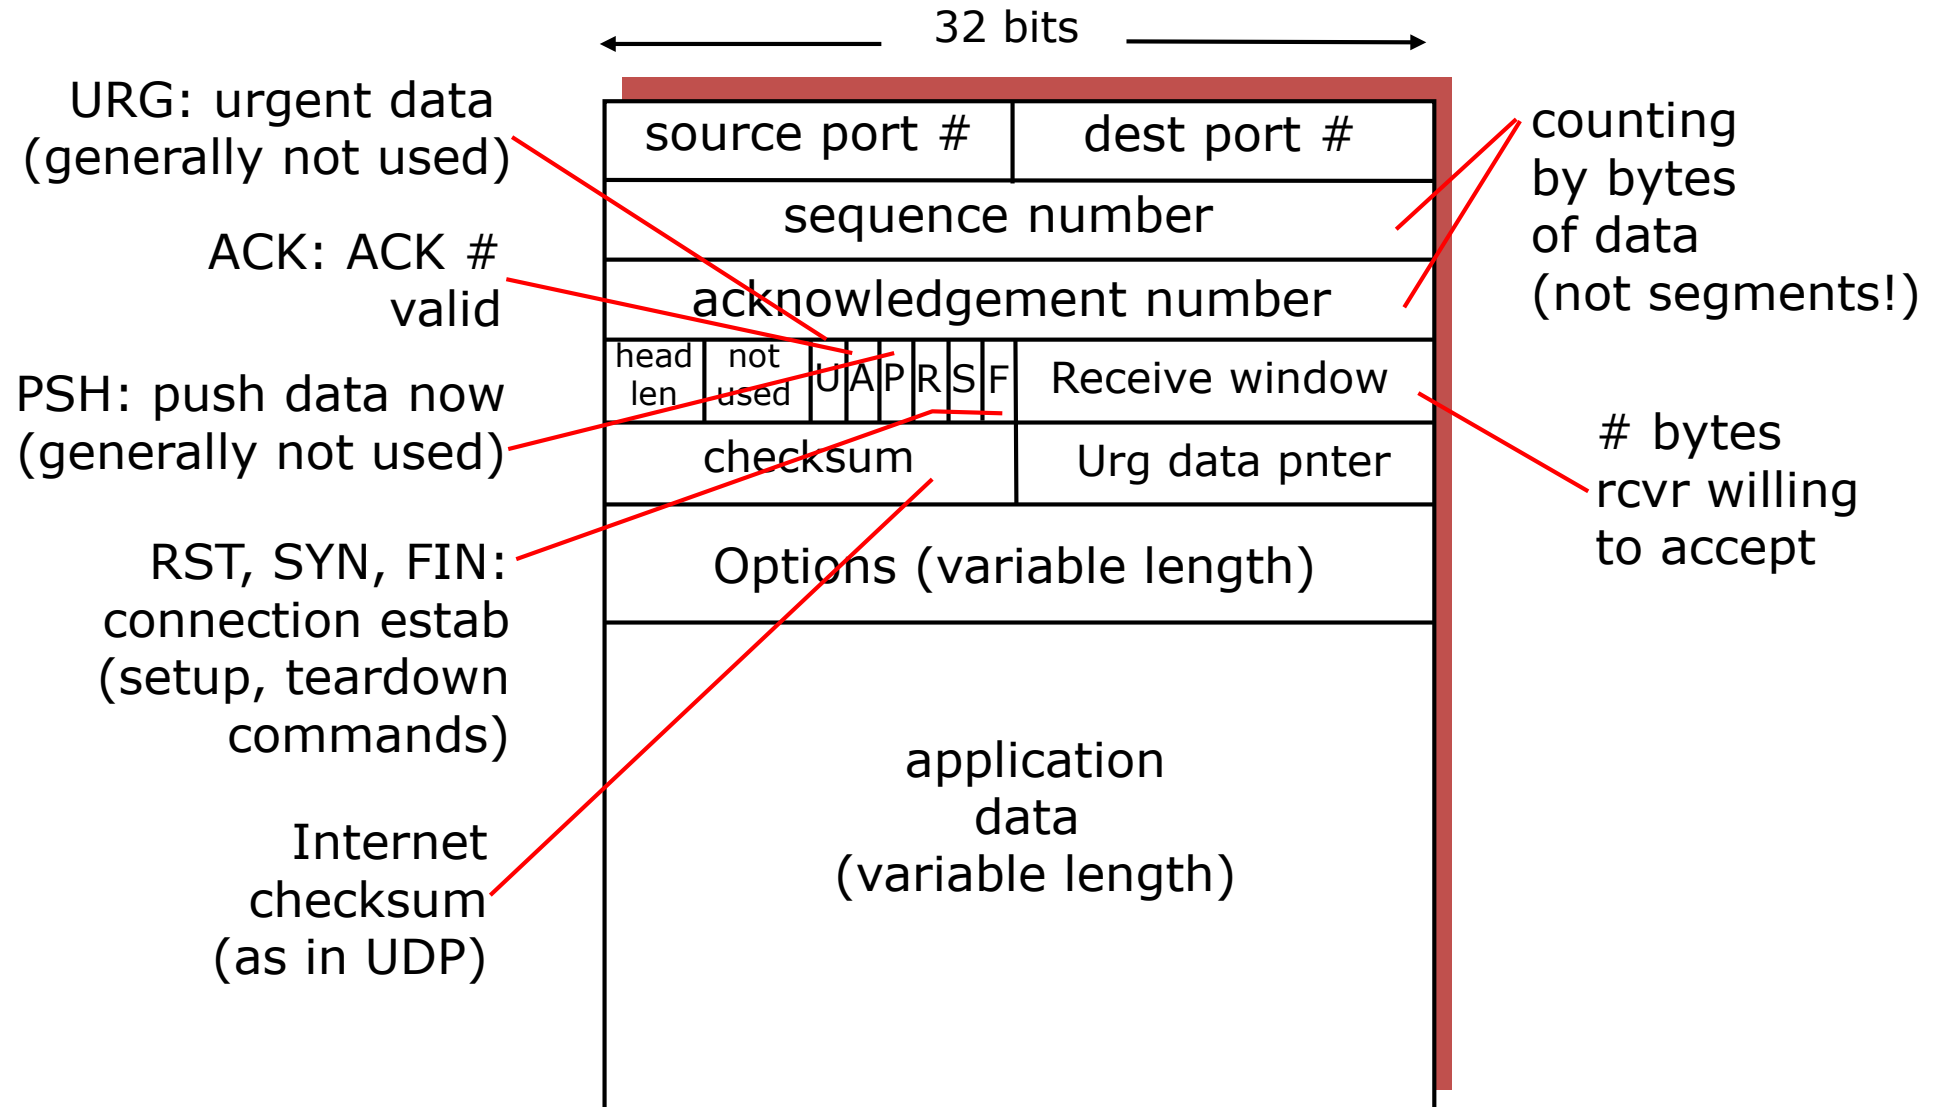
TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

- **point-to-point:**
 - one sender, one receiver
- **reliable, in-order *byte stream*:**
 - no “message boundaries”
- **pipelined:**
 - TCP congestion and flow control set window size
- ***send & receive buffers***
- **full duplex data:**
 - bi-directional data flow in same connection
 - MSS: maximum segment size
- **connection-oriented:**
 - handshaking (exchange of control msgs) init's sender, receiver state before data exchange
- **flow controlled:**
 - sender will not overwhelm receiver

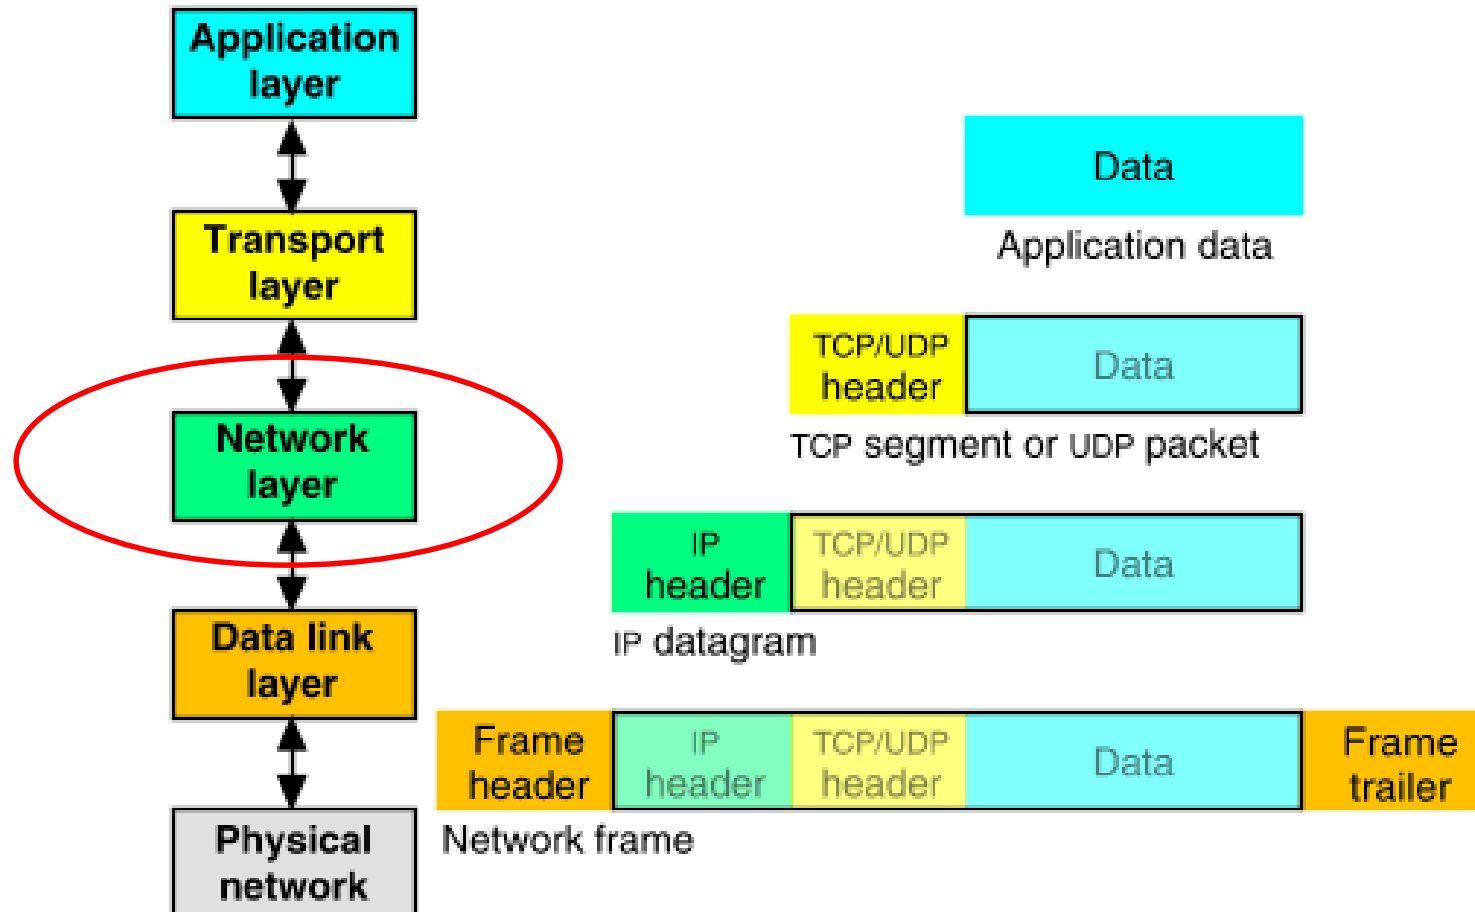


TCP segment structure



THE NETWORK LAYER

IP works at the network layer



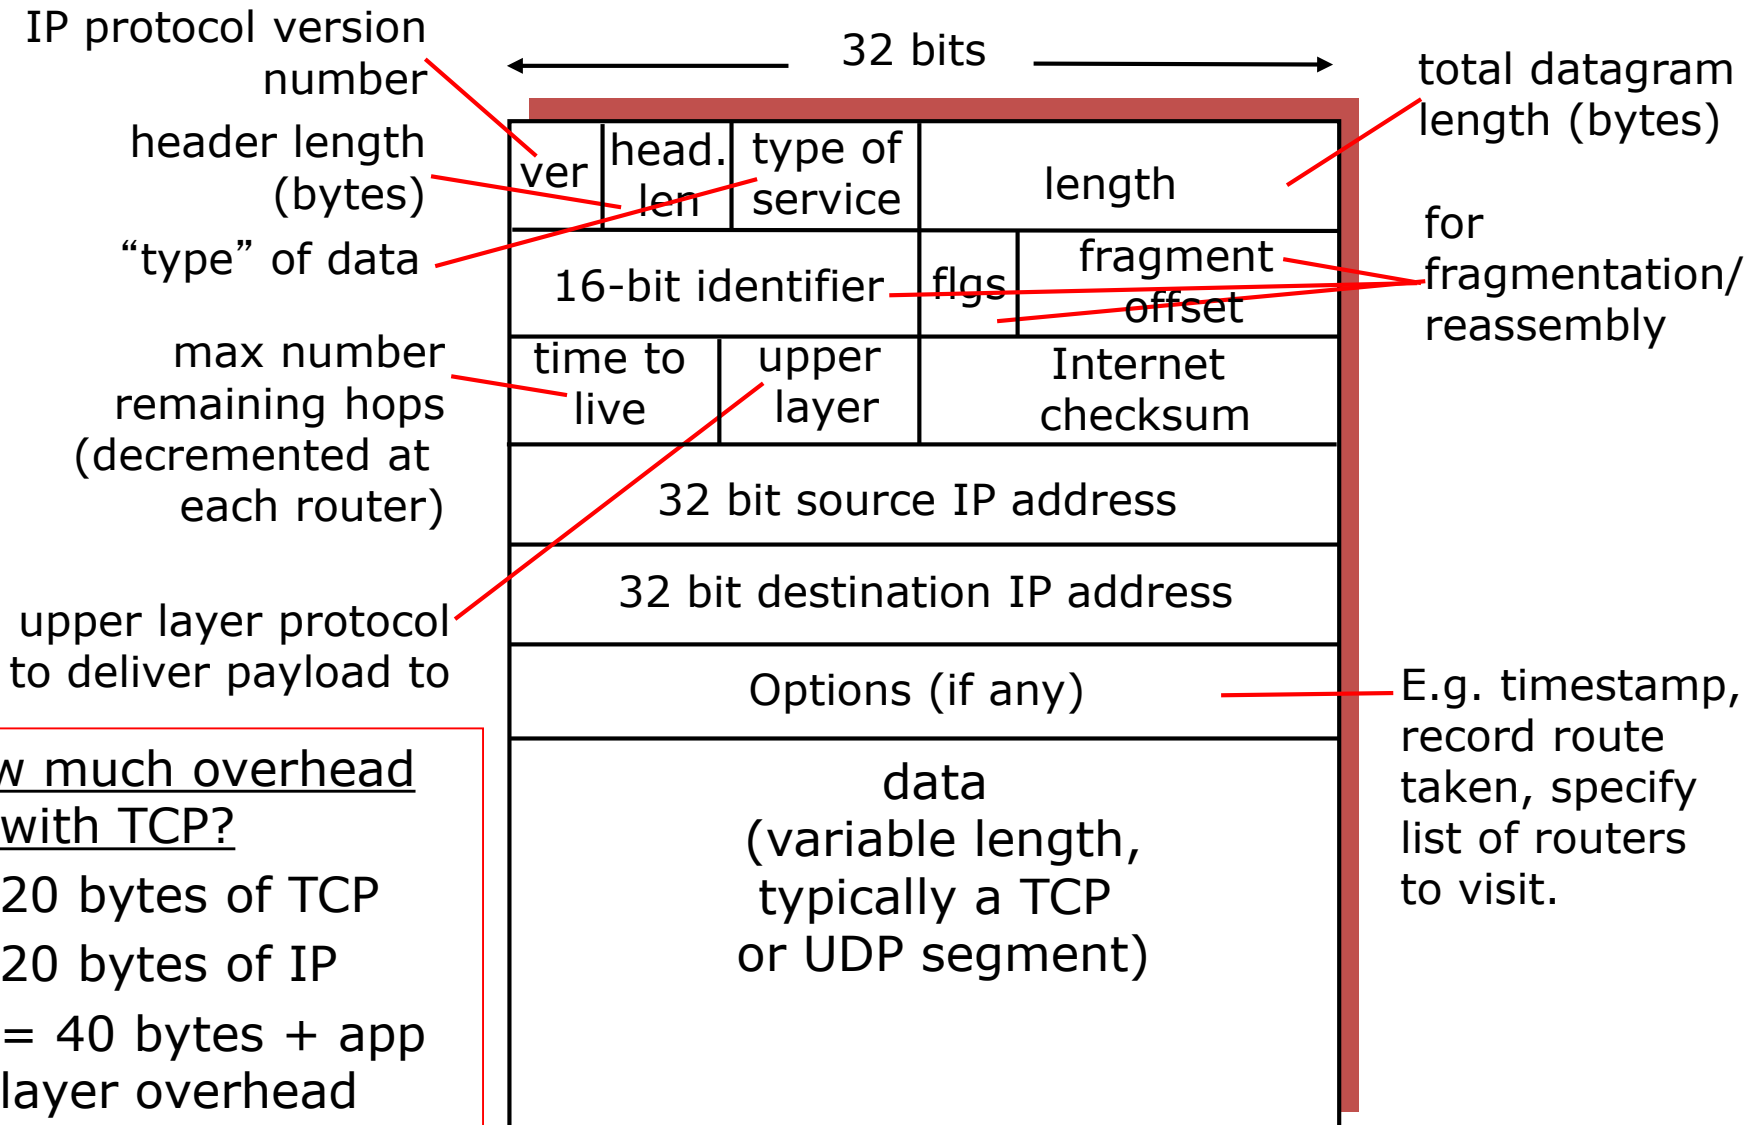
Key Network-Layer Functions

- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to dest.
 - *Routing algorithms*

analogy:

- *routing*: process of planning trip from source to dest
- *forwarding*: process of getting through single interchange

IPv4 datagram format

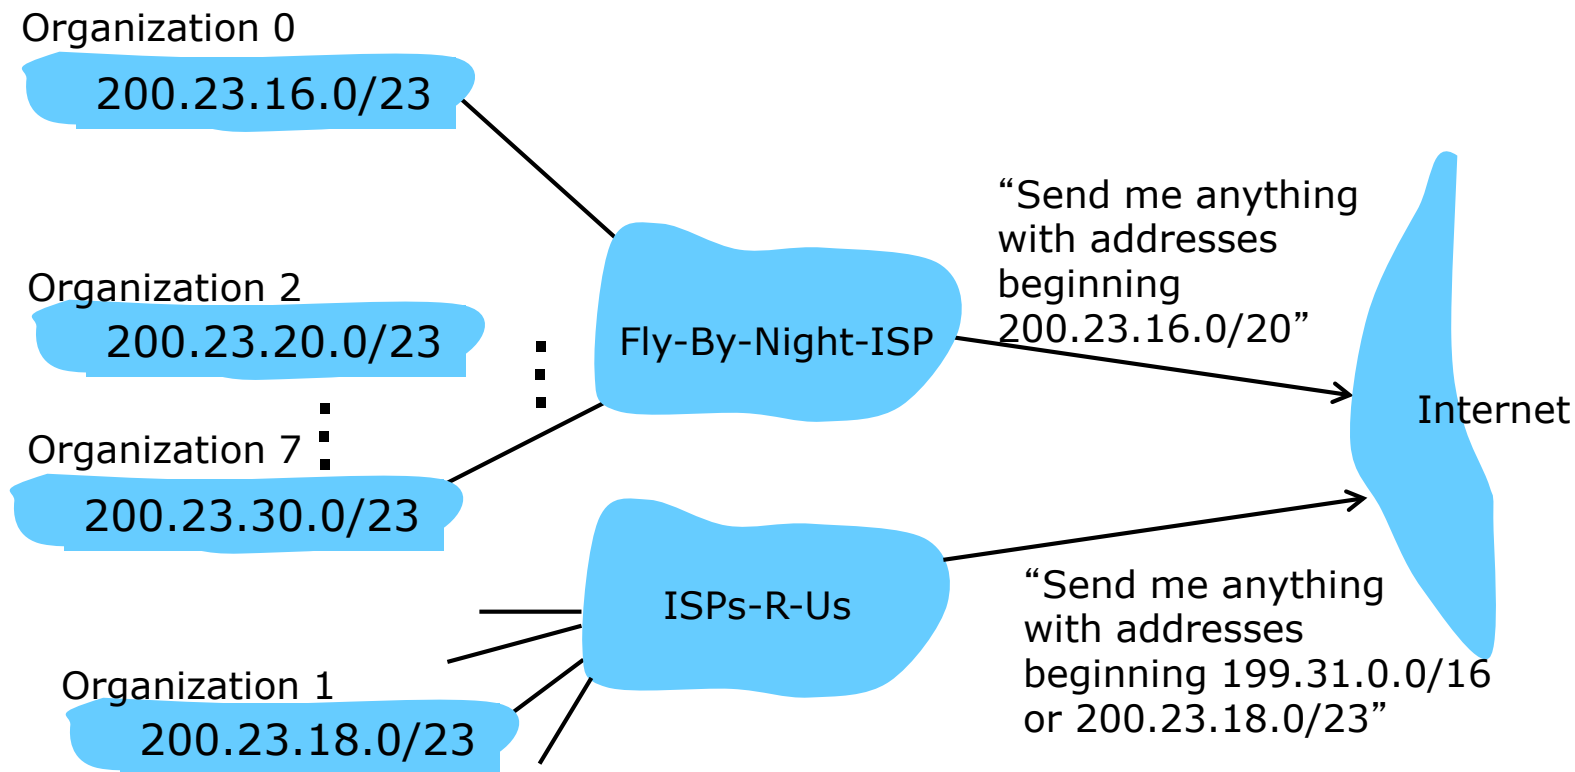


how much overhead with TCP?

- 20 bytes of TCP
- 20 bytes of IP
- = 40 bytes + app layer overhead

Hierarchical addressing: more specific routes

ISPs-R-Us has a more specific route to Organization 1



IP addressing: the last word...

Q: How does an ISP get block of addresses?

A: **ICANN**: Internet **C**orporation for **A**ssigned **N**ames and **N**umbers

- allocates addresses
- manages DNS
- assigns domain names, resolves disputes

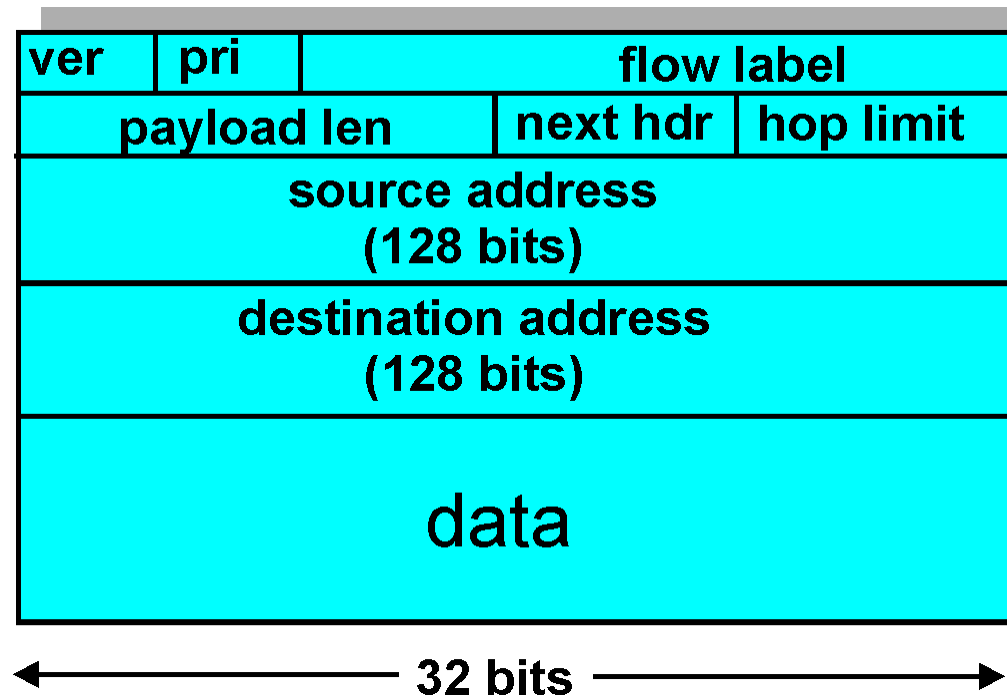
IPv6 datagram format

Priority: identify priority among datagrams in flow

Flow Label: identify datagrams in same “flow.”

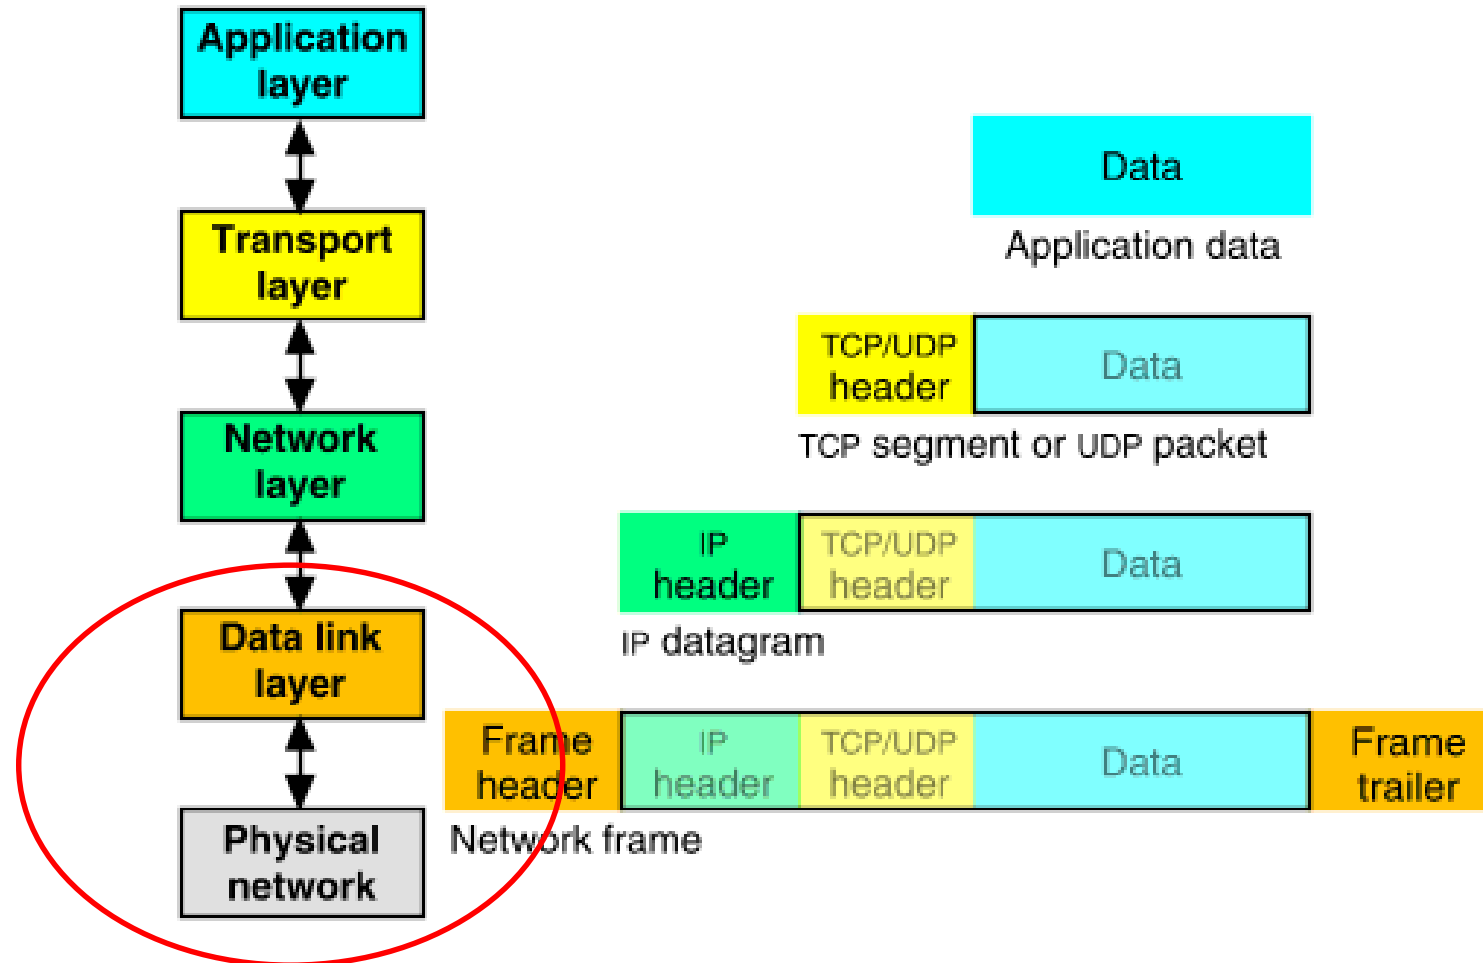
(concept of “flow” not well defined).

Next header: identify upper layer protocol for data



THE LINK AND PHYSICAL LAYERS

Various links and media are possible at the link and physical layers

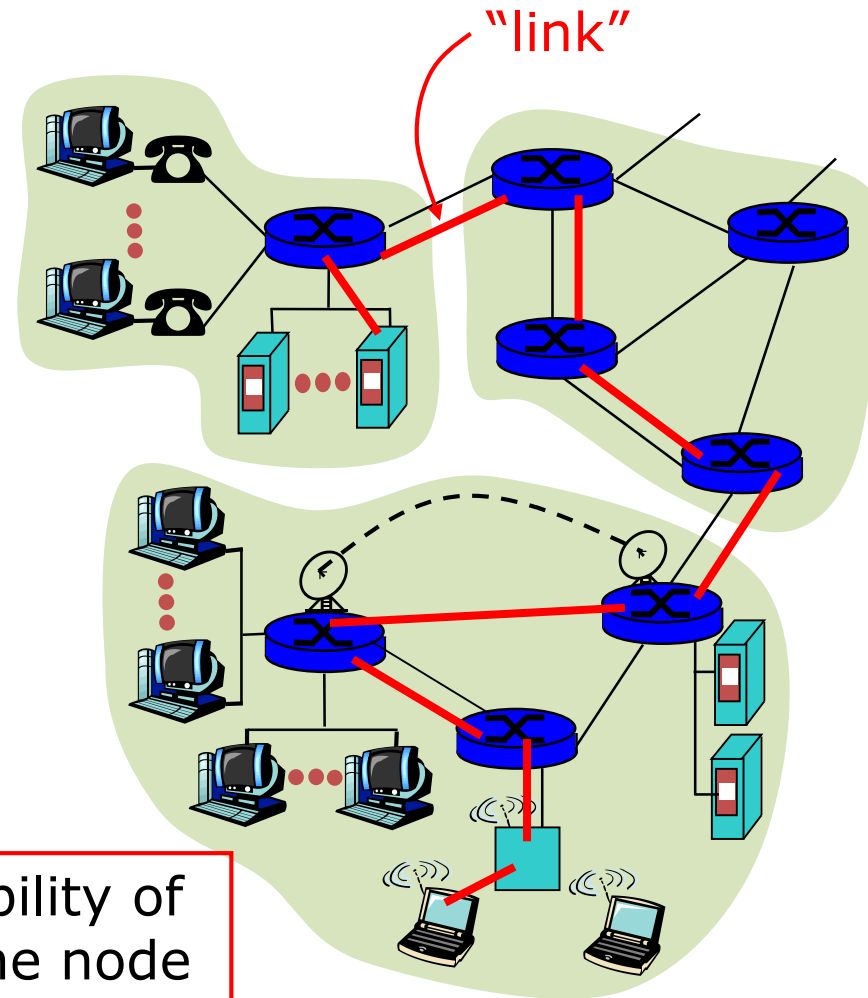


Link Layer: Introduction

Some terminology:

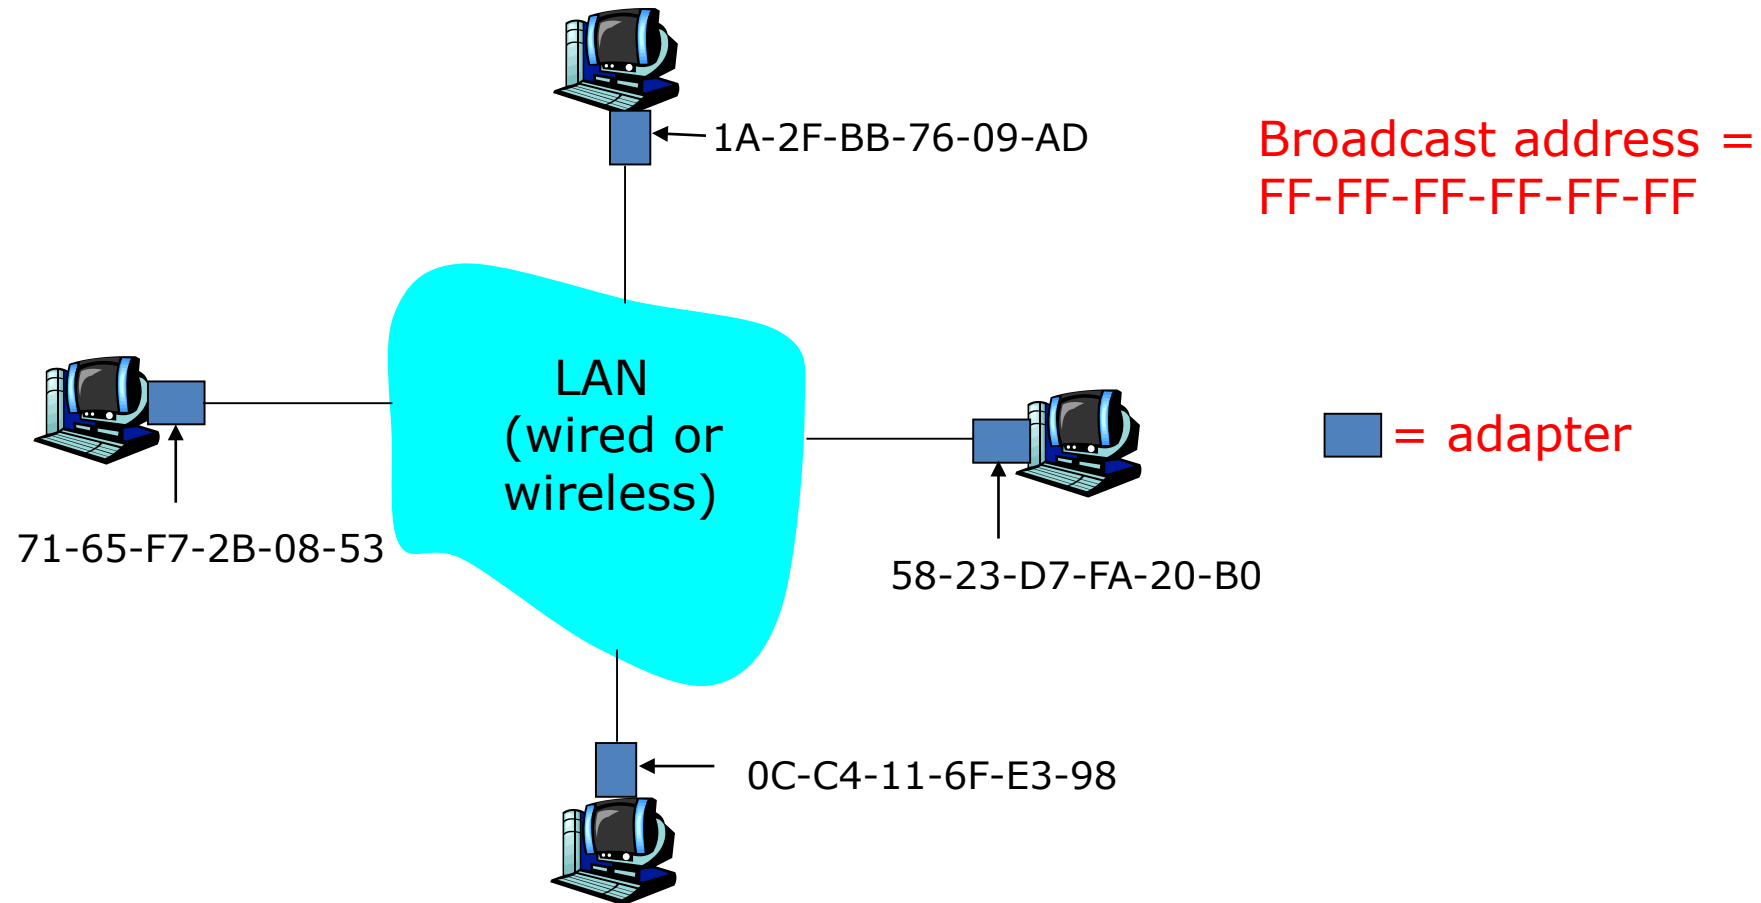
- hosts and routers are **nodes**
- communication channels that connect adjacent nodes along communication path are **links**
 - wired links
 - wireless links
 - LANs
- layer-2 packet is a **frame**, encapsulates datagram

data-link layer has responsibility of transferring datagram from one node to adjacent node over a link



LAN Addresses and ARP

Each adapter on LAN has unique LAN address



Today's Objectives

Brief review of networking

- The layered model
- Application layer
 - HTTP
 - port numbers
- Transport layer
- Network layer
- Link and physical layers