

ECE4574 – Large-Scale SW Development for Engineering Systems

Lecture 13 – Architectural Issues for Engineering Systems

Creed Jones, PhD

Course Updates

- Sprint 2 is underway
- Homework 3 (the final one) will be posted later this week
 - Due November 8
- Quiz 5 is TODAY!!!
 - covers lectures 12-14
 - open 7 PM to 1 AM

Today's Objectives

Some example architectural patterns

- N-tier Client Server
- Messaging
- Publish-Subscribe
- Broker
- Process Coordinator



Embedded Systems Software Architectures

- Software Architecture for Embedded Systems
- Automotive Software Engineering

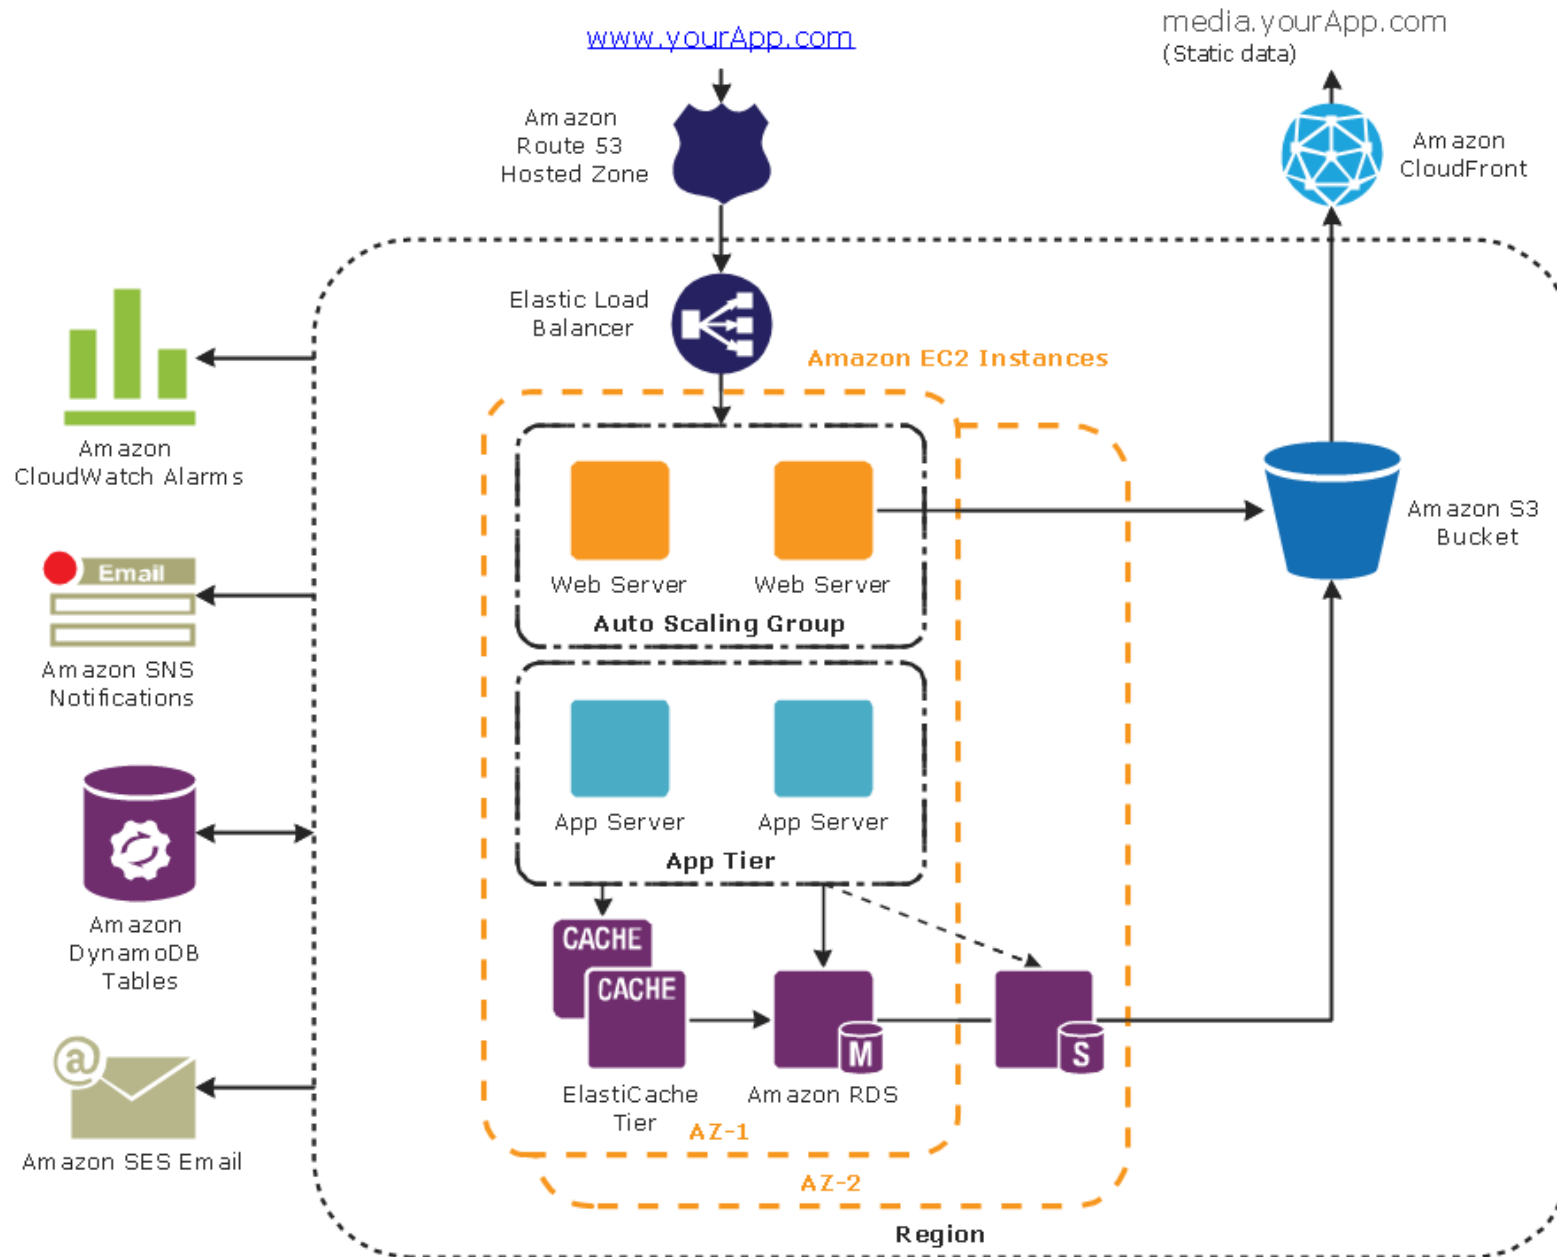
ARCHITECTURAL PATTERNS

There is a distinction between an *Architectural Pattern* and a *Design Pattern*

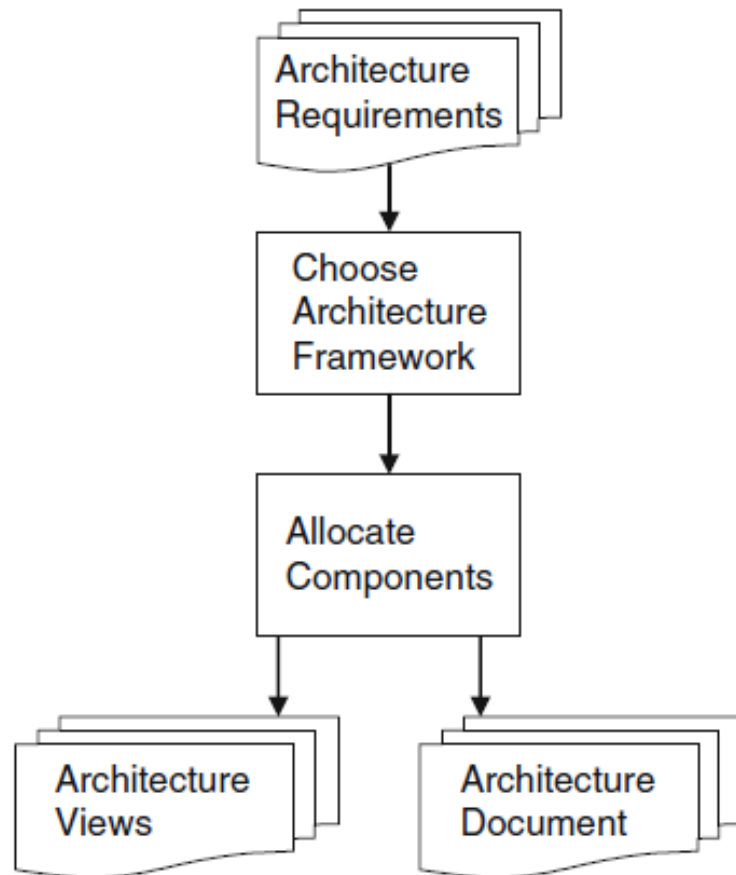
- The architectural pattern is a time-tested arrangement of overall components that defines the structure of the entire project
- Design patterns are used to specify the shape and working of specific pieces
- In classical church architecture, the basilica is an Architectural Pattern
- A vaulted arch is a design pattern



3-Tier Auto-scalable Web Application Architecture



"The design stage itself has two steps...choosing an overall strategy for the architecture, based around proven architecture patterns...[and] specifying the individual components that make up the application"



- Note that this process does not iterate
 - hopefully
- We choose a framework from past experience and familiarity, aligned with the requirements
- That leads to defining components

ARCHITECTURAL PATTERNS – SOME EXAMPLES

The N-tier client server framework supports web applications

1. Separation of Concerns

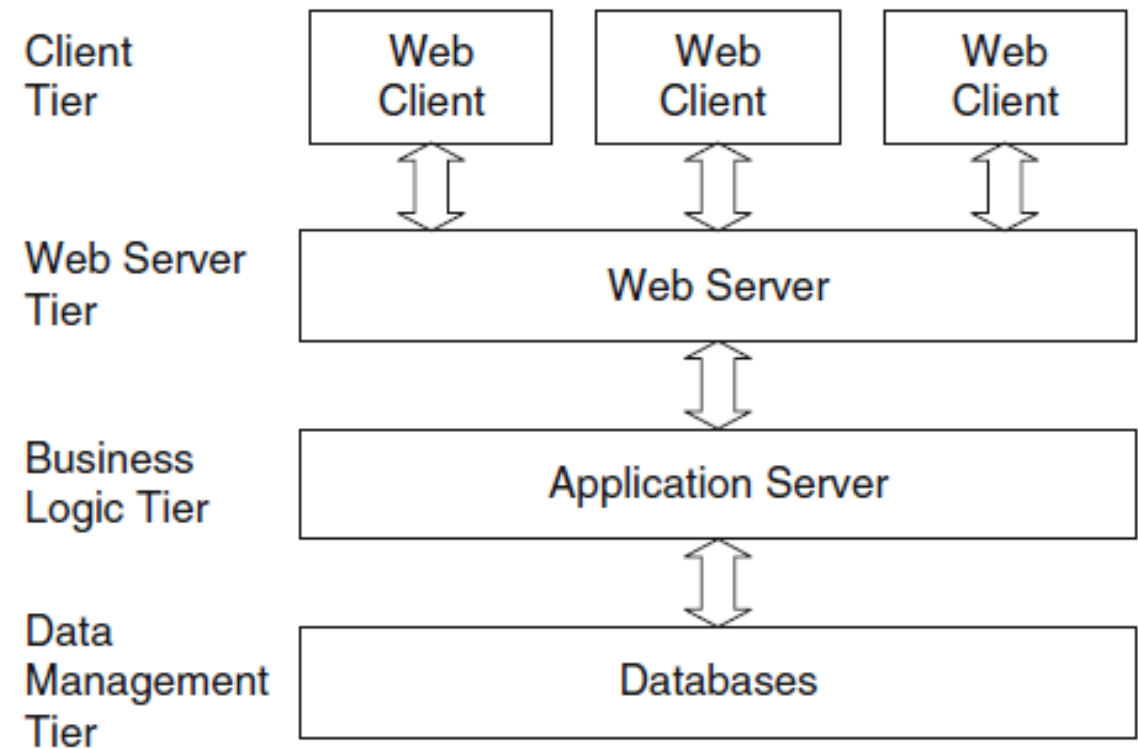
- Presentation, business and data are handled in different tiers

2. Synchronous Communications

- Tiers talk to each other via "request-reply"
- Tiers wait for a response before proceeding

3. Flexible Deployment


- One or multiple machines



How does the N-Tier Client Server meet general quality requirements?

| Quality attribute | Issues |
|-------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| Availability | Servers in each tier can be replicated for reliability (at degraded performance) |
| Failure handling | If a server fails while a client is communicating, we want the request routed to a live server for completion (most app servers do this) |
| Modifiability | Most changes will be limited to one tier of the design |
| Performance | High performance, if the servers support many concurrent threads and the connection between tiers is lightweight |
| Scalability | Servers can be replicated up to a point; the data management tier will often become the bottleneck |

The Messaging pattern is a simple concept in which all components communicate by FIFO queues with non-trivial size




1. Asynchronous Communication
 - Clients generally do not block on responses; messages are handled by the queue until a server accepts it
2. Configurable Quality of Service (QoS)
 - May be unreliable or reliable, slow or fast
3. Loose Coupling
 - Servers process all messages; clients don't send requests to specific servers (sources and destinations are not known)

The Messaging architecture pattern has specific quality characteristics

| Quality attribute | Issues |
|-------------------|------------------------------------------------------------------------------------------------------|
| Availability | Duplicate queues with the same name on different machines provide redundancy |
| Failure handling | If a queue fails while a client is communicating, it can find a replica queue and repost |
| Modifiability | If messages are self-describing (XML), changes should be localized |
| Performance | Messaging can have very high throughput, depending on the reliability requirement |
| Scalability | Queues can reside on the endpoints or dedicated messaging servers; this is a highly scalable pattern |

The Publish-Subscribe architectural pattern allows for one message to be received by many destination processes




1. Many-to-many messaging
 - Messages are received by any subscriber who "registers" for it, by *topic*
2. Configurable QoS
 - Reliable or unreliable, but also point-to-point or multicast
3. Loose Coupling
 - Servers process all messages; clients don't send requests to specific servers (sources and destinations are not known)

Publish-subscribe has good quality in many areas of concern

| Quality attribute | Issues |
|-------------------|------------------------------------------------------------------------------------------------------------------------------|
| Availability | Duplicate queues for the same topic on different machines provide redundancy |
| Failure handling | If a topic server fails while a client is communicating, it can find a replica server and repost |
| Modifiability | Loose coupling allows modification |
| Performance | Publish-subscribe can have very high throughput, depending on the reliability requirement (especially for multicast publish) |
| Scalability | Topics can be hosted on independent servers or clusters; this is a highly scalable pattern |

The Broker pattern allows for exchange of messages that demand translation or other processing

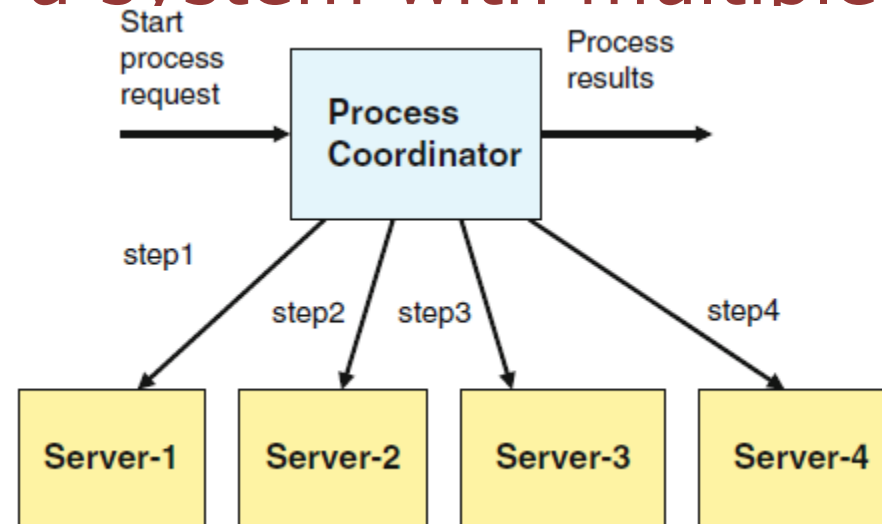


1. Hub-and-spoke architecture
 - the broker is a central messaging hub; ports are associated with a specific message format
2. Performs message routing
 - Logic to determine message paths is in the router
3. Performs message transformation
 - Message type to message type, as determined by the source and destination port

The Broker pattern and its quality characteristics

| Quality attribute | Issues |
|-------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| Availability | High availability will require that the broker be replicated |
| Failure handling | Input ports expect particular message types; incompatible messages are discarded; senders can find a live broker in the case one fails |
| Modifiability | All changes in format or sequence are handled in the broker |
| Performance | The broker component can become a bottleneck |
| Scalability | Implementing the broker on a cluster of machines allows the Broker pattern to scale |

The Process Coordinator can implement complex business rules in a system with multiple servers



1. Process encapsulation
 - Sequence of steps needed to accomplish the business process; receives a request and responds by initiating a list of steps
2. Loose coupling
 - Servers are not aware of the larger business context, but merely respond to requests
3. Flexible communications
 - Each path can be synchronous or asynchronous (queued); may use a callback function

A Process Coordinator pattern has some good quality attributes and some that are so-so

| Quality attribute | Issues |
|-------------------|--------------------------------------------------------------------------------------------------------|
| Availability | We must replicate the coordinator to avoid a single point of failure |
| Failure handling | Lengthy business sequences dictate difficult and complex failure recovery mechanisms |
| Modifiability | The process coordinator holds the definition of the business process in one place |
| Performance | The coordinator must be able to handle multiple requests, perhaps at different points in their process |
| Scalability | We can replicate the coordinator to load-balance and allow scaling |

Once an architectural framework is chosen, we allocate components (what are the large pieces of the project)

- Identifying the major application components, and how they plug into the framework.
- Identifying the interface or services that each component supports.
- Identifying the responsibilities of the component, stating what it can be relied upon to do when it receives a request.
- Identifying dependencies between components.
- Identifying partitions in the architecture that are candidates for distribution over servers in a network.

When identifying components, we remember our usual good design principles

- Minimize dependencies between components (loose coupling)
- Highly cohesive components
- Keep dependencies on COTS components (like middleware, DB, etc) localized
- Structure components hierarchically
- Minimize calls between components
- Look for places to apply design patterns

EMBEDDED SYSTEMS SOFTWARE ARCHITECTURES

Software Engineering for Embedded Systems


Software Engineering for Embedded Systems; Methods, Practical Techniques, and Applications, Newnes, 2013

ISBN 978-0-12-415917-4

Available through the VT library

This book highlights three important points for embedded systems architecture:

1. Constraints are paramount
2. Some new design patterns are useful
3. Deployment must be considered



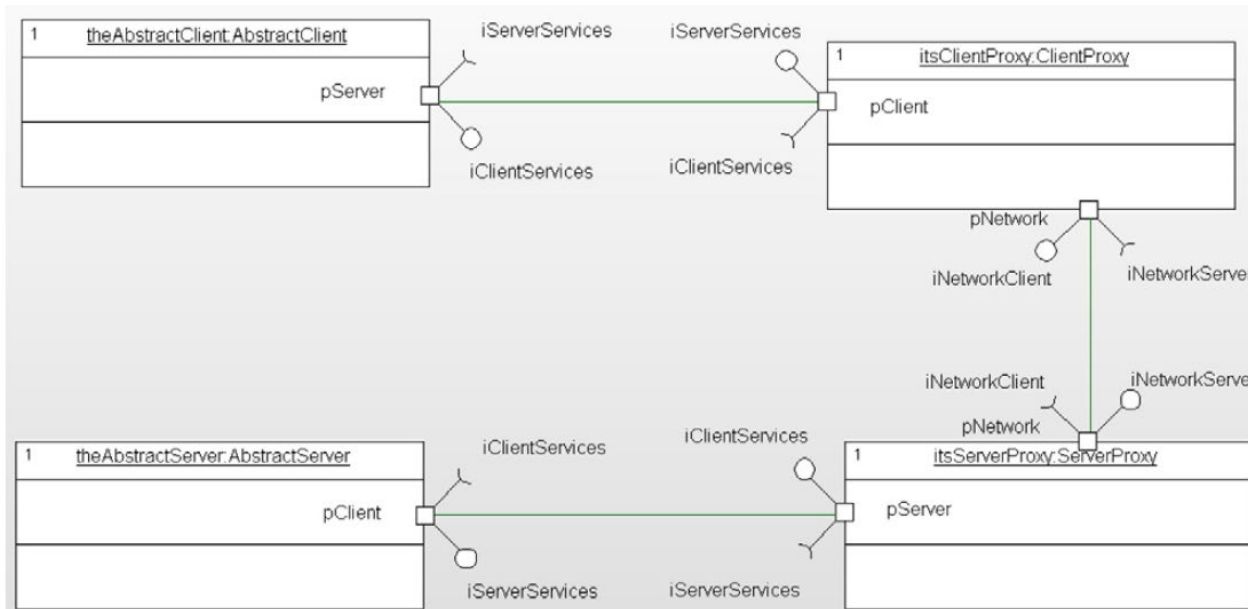
“Architecture is about system-wide optimization”

- Embedded systems usually have strict performance constraints
 - Throughput, reaction time, accuracy, resources, uptime...
- Hard vs. Soft real-time systems
 - Soft: faster response is preferred, and occasional reaction slightly over the “deadline” can be acceptable
 - Hard: deadlines are absolute
- Failure recovery is crucial
 - There’s no one there to reboot

Some new or enhanced design patterns are useful in embedded systems

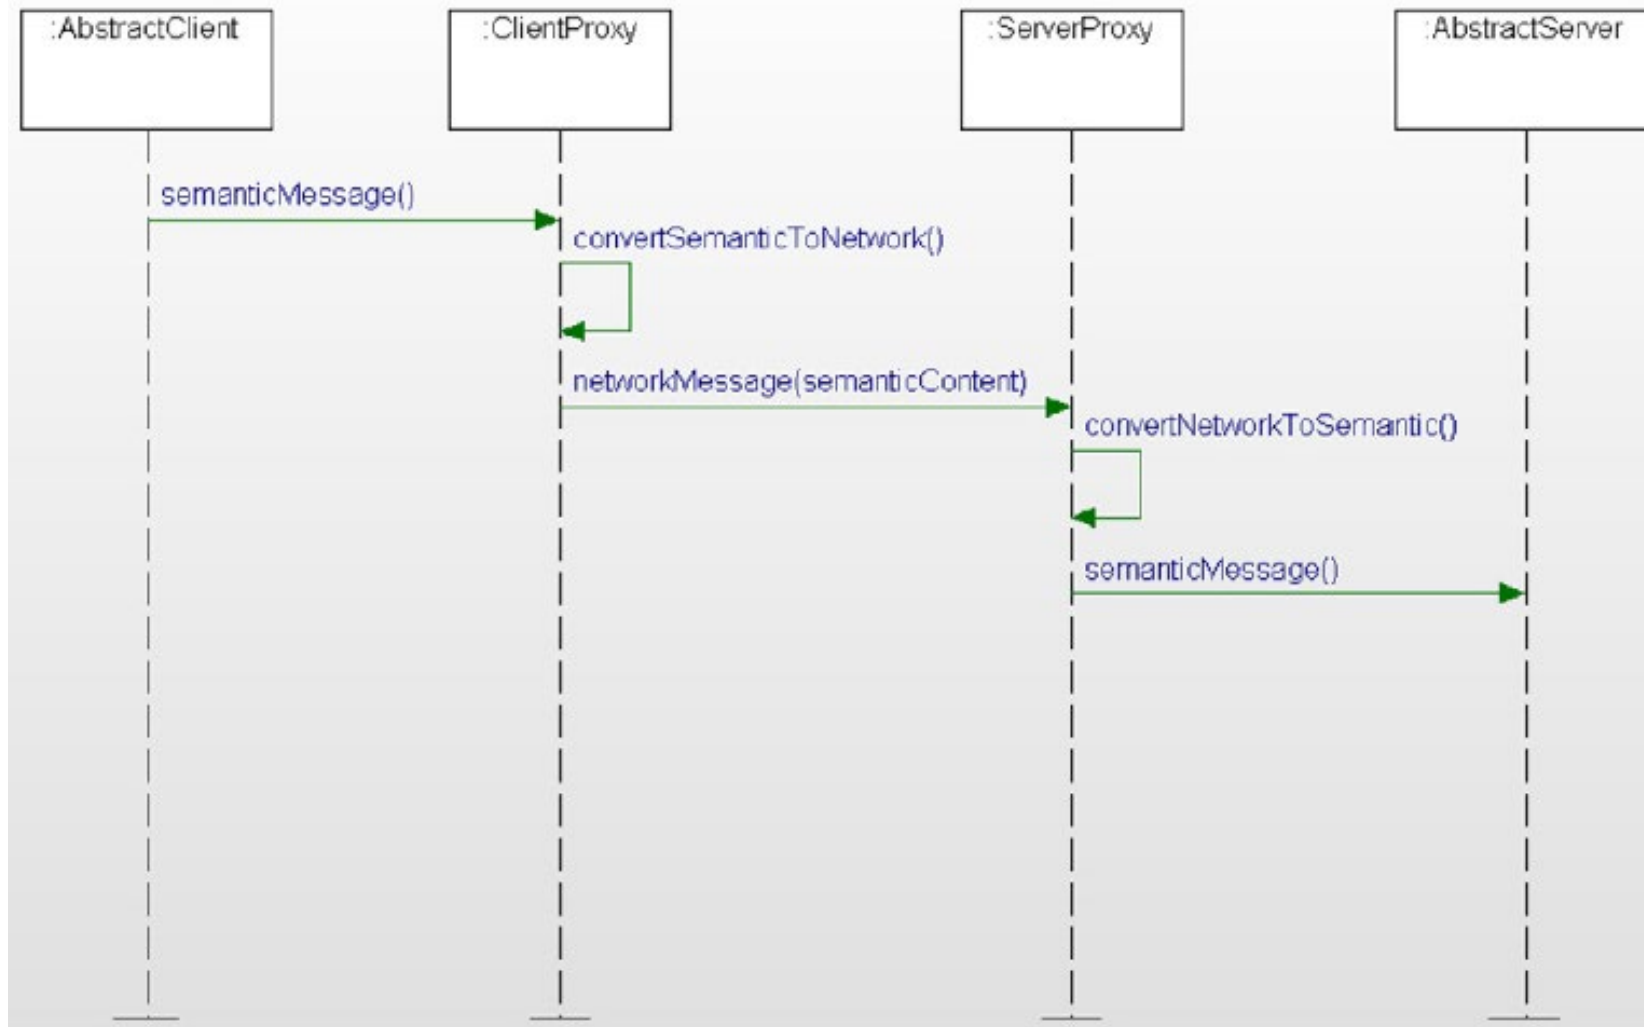
- Port proxy pattern
- Scheduling patterns:
 - Cyclic executive
 - Time-triggered cyclic executive
 - Rate monotonic scheduling (RMS)
 - Earliest deadline first (EDF)

In the *port proxy* pattern, a number of units are connected via discrete connections called ports



- The client and the server are both coded without knowledge of the messaging channel details
- Each has a proxy that handles conversion to/from the protocol of the channel
- Changing to a new network or middleware would only require updates to the proxies

The *semantic message* (the essence of the message) is converted to and from the network format

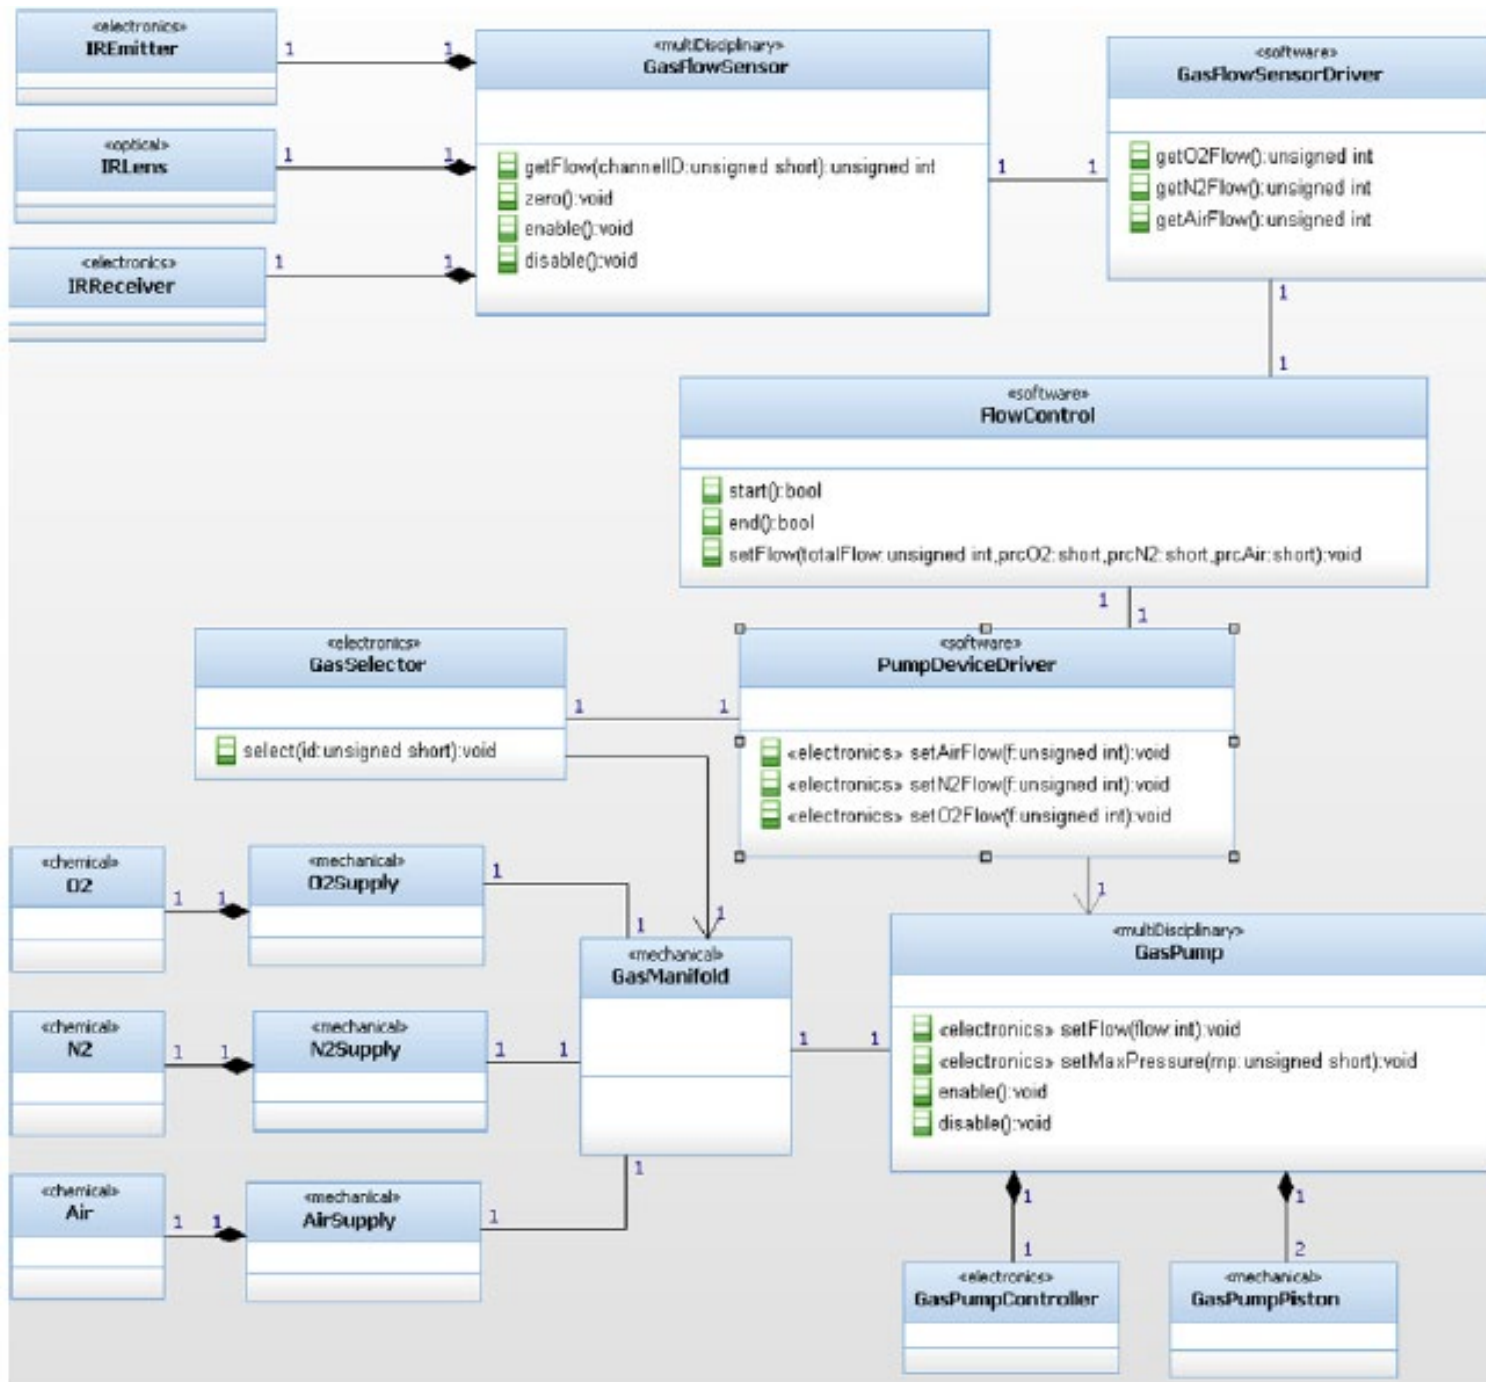


The various *scheduling* patterns enable implementations of different ways of sequencing and assigning tasks

| Pattern | Description | Benefits | Drawbacks |
|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|
| Cyclic executive | The scheduler runs a list of tasks (each to completion) in the same order in a repetitive cycle. | Simple Fair Highly predictable | Low responsiveness Unstable Suboptimal performance Requires tuning |
| Time-triggered cyclic executive | Same as cyclic executive except that each cycle begins on a time-based epoch. | Simple Fair Highly predictable Synchronizes with reference clock | Low responsiveness Unstable Suboptimal performance Requires tuning |
| Rate monotonic scheduling (RMS) | All tasks are assumed to be periodic with the deadline at the end of the period. Priorities are assigned at design time on the basis of period - shorter periods have higher priority. Highest priority task always runs. | Stable Optimal Robust | Unfair May not scale to very complex systems More complex Less predictable |
| Earliest deadline first (EDF) | Priorities are assigned at run-time based on the nearness of the deadline (i.e., its urgency). Highest priority waiting task always runs. | Optimal Robust | Unfair Naïve implementation Leads to thrashing Unstable More complex Less predictable |

The deployment of software is a different and important concern for embedded systems design

- What code will run where?
- What “OS” facilities and platforms will be available?
- Sometimes, the connection between system elements is:
 - Sporadic: think of cell connection in outlying areas
 - Occasional: scheduled connection to hosts (overnight?)
- Often, code deploys other code
- Bootloading and startup must be supported by system architecture



Note that the deployment diagram shows sensors and mechanical elements (!)


Automotive Software Engineering

Schaeuffele, J., & Zurawka, T.
(2016). *Automotive software engineering, second edition*. ProQuest Ebook
Central <https://ebookcentral.proquest.com>

Available through the VT library

Applies software engineering principles to automotive systems, which are:

- real-time
- modular
- networked
- cost-sensitive



CAN is a multi-master serial messaging bus; FlexRay is faster and more reliable but more expensive

On occasion, Ethernet is also used for automotive systems

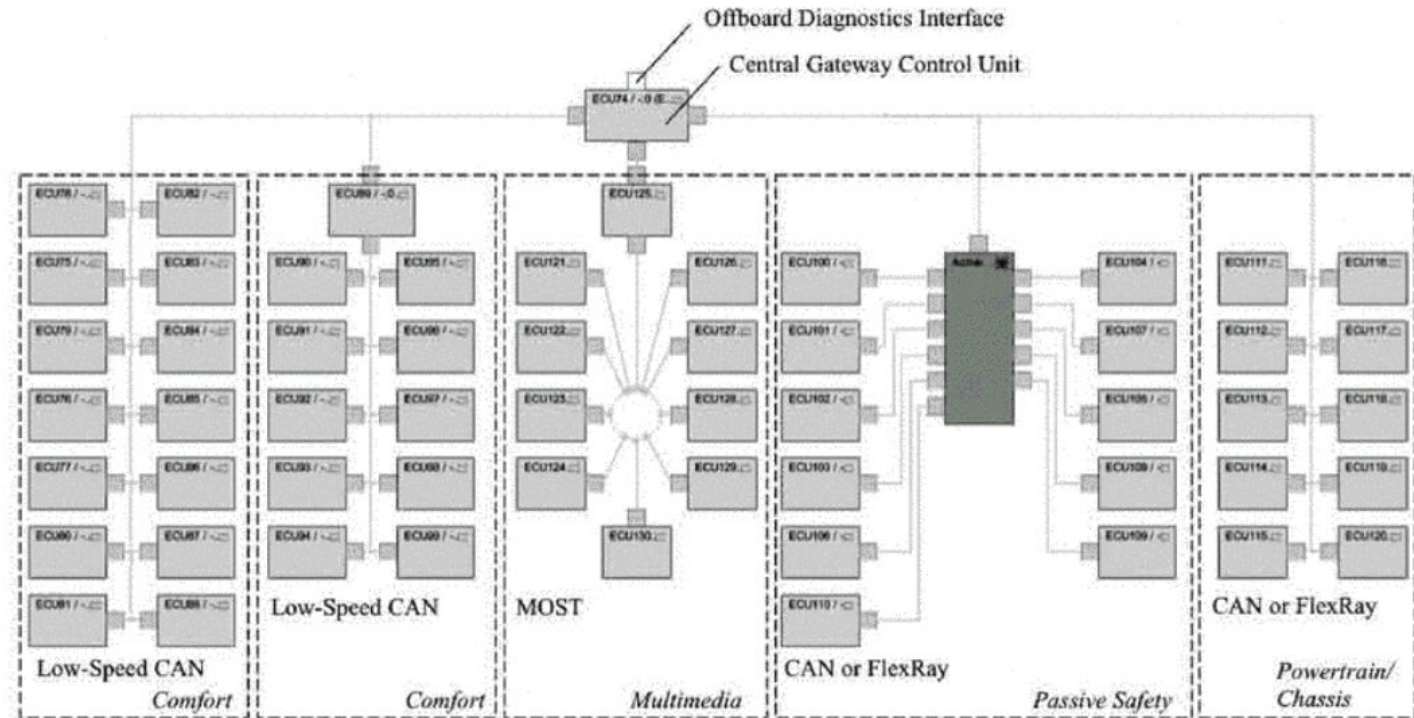


Figure 2.66 ECU Network of a premium class vehicle.

Automotive systems software is highly modular and component-based

- Software components mirror hardware components
- Some real-time events are tolerant of response lag
 - comfort systems
- Others are not
 - engine control and safety systems

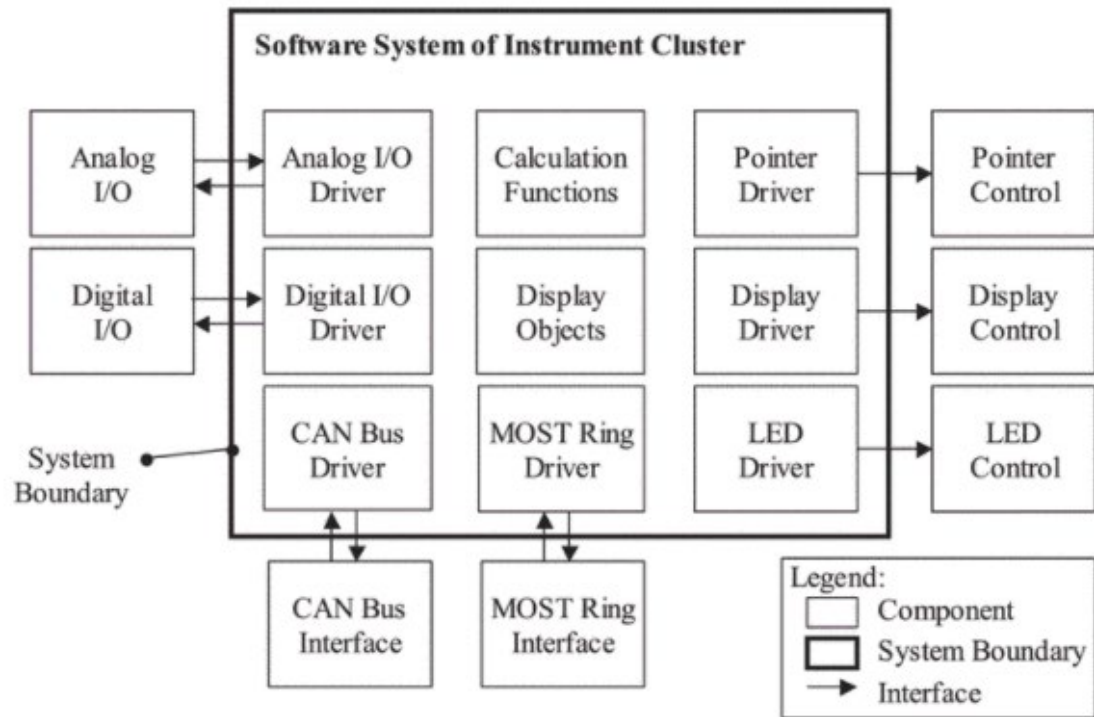


Figure 4.20 Context and interface model of the instrument cluster software.

AUTOSAR is a consortium-based standard for open architecture for automotive systems




Figure 4.22 Software architecture of the instrument cluster.

- https://www.autosar.org/fileadmin/ABOUT/AUTOSAR_EXP_Introduction102020.pdf
- It defines three layers of software functions:
 - basic software: hardware support and general functions
 - runtime environment: messaging
 - application software: specific to the automotive task
- The AUTOSAR runtime serves as component middleware

Specification of automotive SW components reflects the needs of the application

- The data model for information to be processed:
 - scalar, vector, matrix
- The behavioral model (control flow)
- The real-time model (latency and throughput guarantees)
- Note the importance of testing

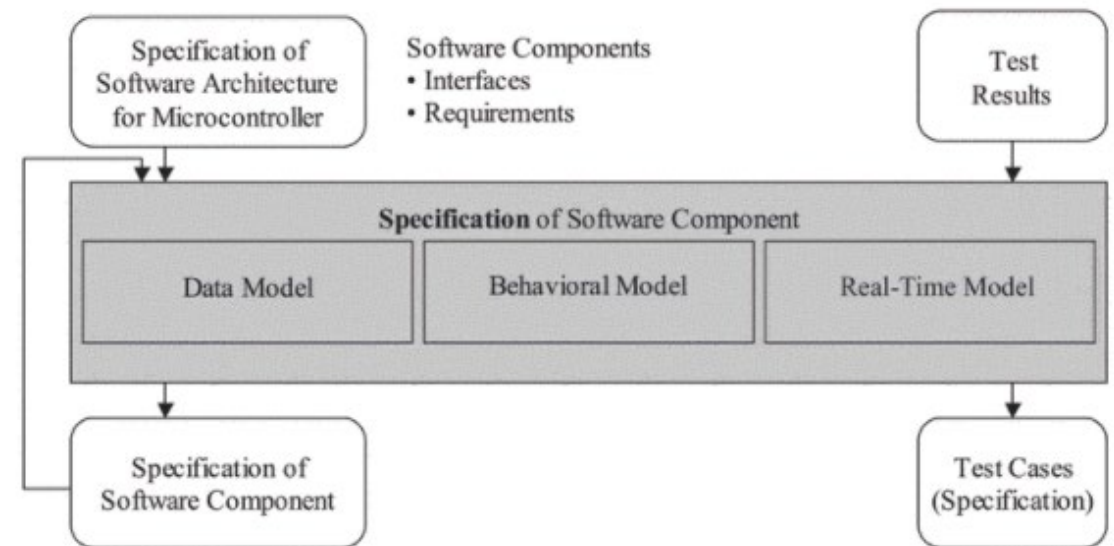


Figure 4.24 Specification of a software component.

Today's Objectives

Some example architectural patterns

- N-tier Client Server
- Messaging
- Publish-Subscribe
- Broker
- Process Coordinator

Embedded Systems Software Architectures

- Software Architecture for Embedded Systems
- Automotive Software Engineering